

SkePU

<https://skepu.github.io>

SkePU Framework Introduction & Tutorial

August Ernstsson Christoph Kessler

{firstname.lastname}@liu.se

Linköping University, Sweden

Basic Topics

- Skeleton Programming
- SkePU Introduction and History
- Using SkePU
- Compilation Options and Backend Selection
- Smart Containers and Consistency Model

Advanced Topics

- SkePU in Current Research
 - Lazy Skeleton Evaluation
 - Multi-Variant User Functions
 - Deterministic Pseudo-Random Number Generation


Intermediate Topics

- Skeletons in Depth
 - Map
 - Reduce
 - MapReduce
 - Scan
 - MapOverlap
 - MapPairs
 - MapPairsReduce
- Nested User Functions
- SkePU Standard Library
 - BLAS
 - Deterministic PRNG
 - Image Filters
 - Utilities

Other

- Live Demo
- Behind the Scenes

- Currently, the best information resource on SkePU is August Ernstssons doctoral thesis (2022)
- <https://doi.org/10.3384/9789179291969>
- See also the user guide on <https://skepu.github.io> for more concrete instructions on e.g. installation.
- Revised user guide and documentation is upcoming!



Skeleton Programming

Programming parallel systems is hard!

- Resource utilization
- Synchronization, Communication
- Memory consistency
- Different hardware architectures, heterogeneity


Skeleton programming (algorithmic skeletons)

- A high-level parallel programming concept
- Inspired by functional programming
- Generic computational patterns
- Abstracts architecture-specific issues

Skeletons


Parametrizable higher-order constructs

- Map
- Reduce
- MapReduce
- Scan
- and others




User functions

User-defined operators




Skeleton parametrization example

Dot product operation


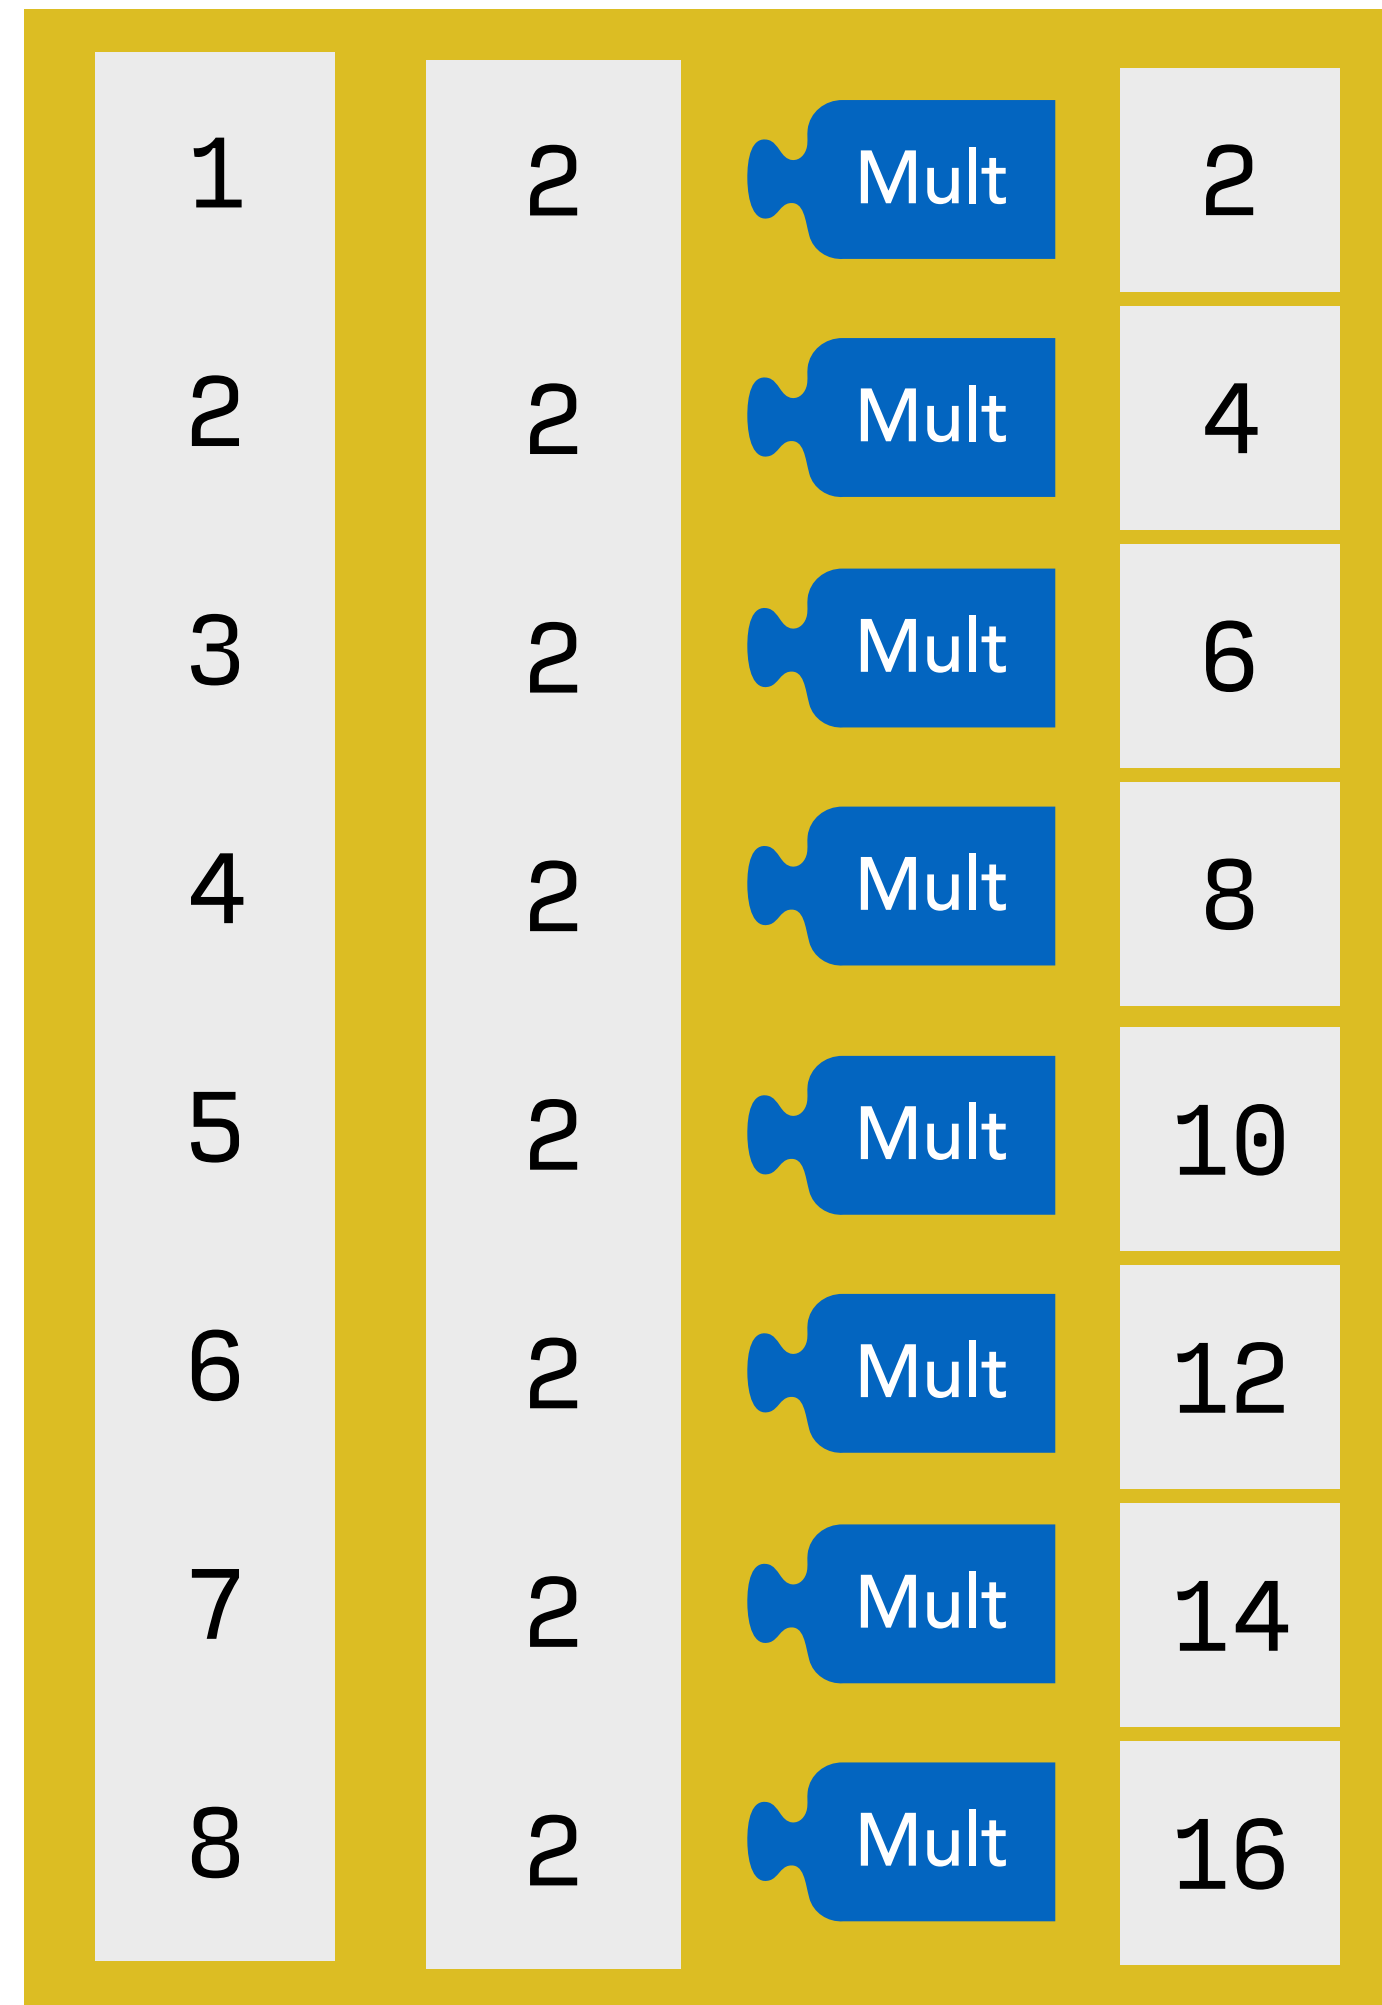


Sequential algorithm


1	2	Mult	2
2	2	Mult	4
3	2	Mult	6
4	2	Mult	8
5	2	Mult	10
6	2	Mult	12
7	2	Mult	14
8	2	Mult	16








Parallel algorithm



Sequential algorithm



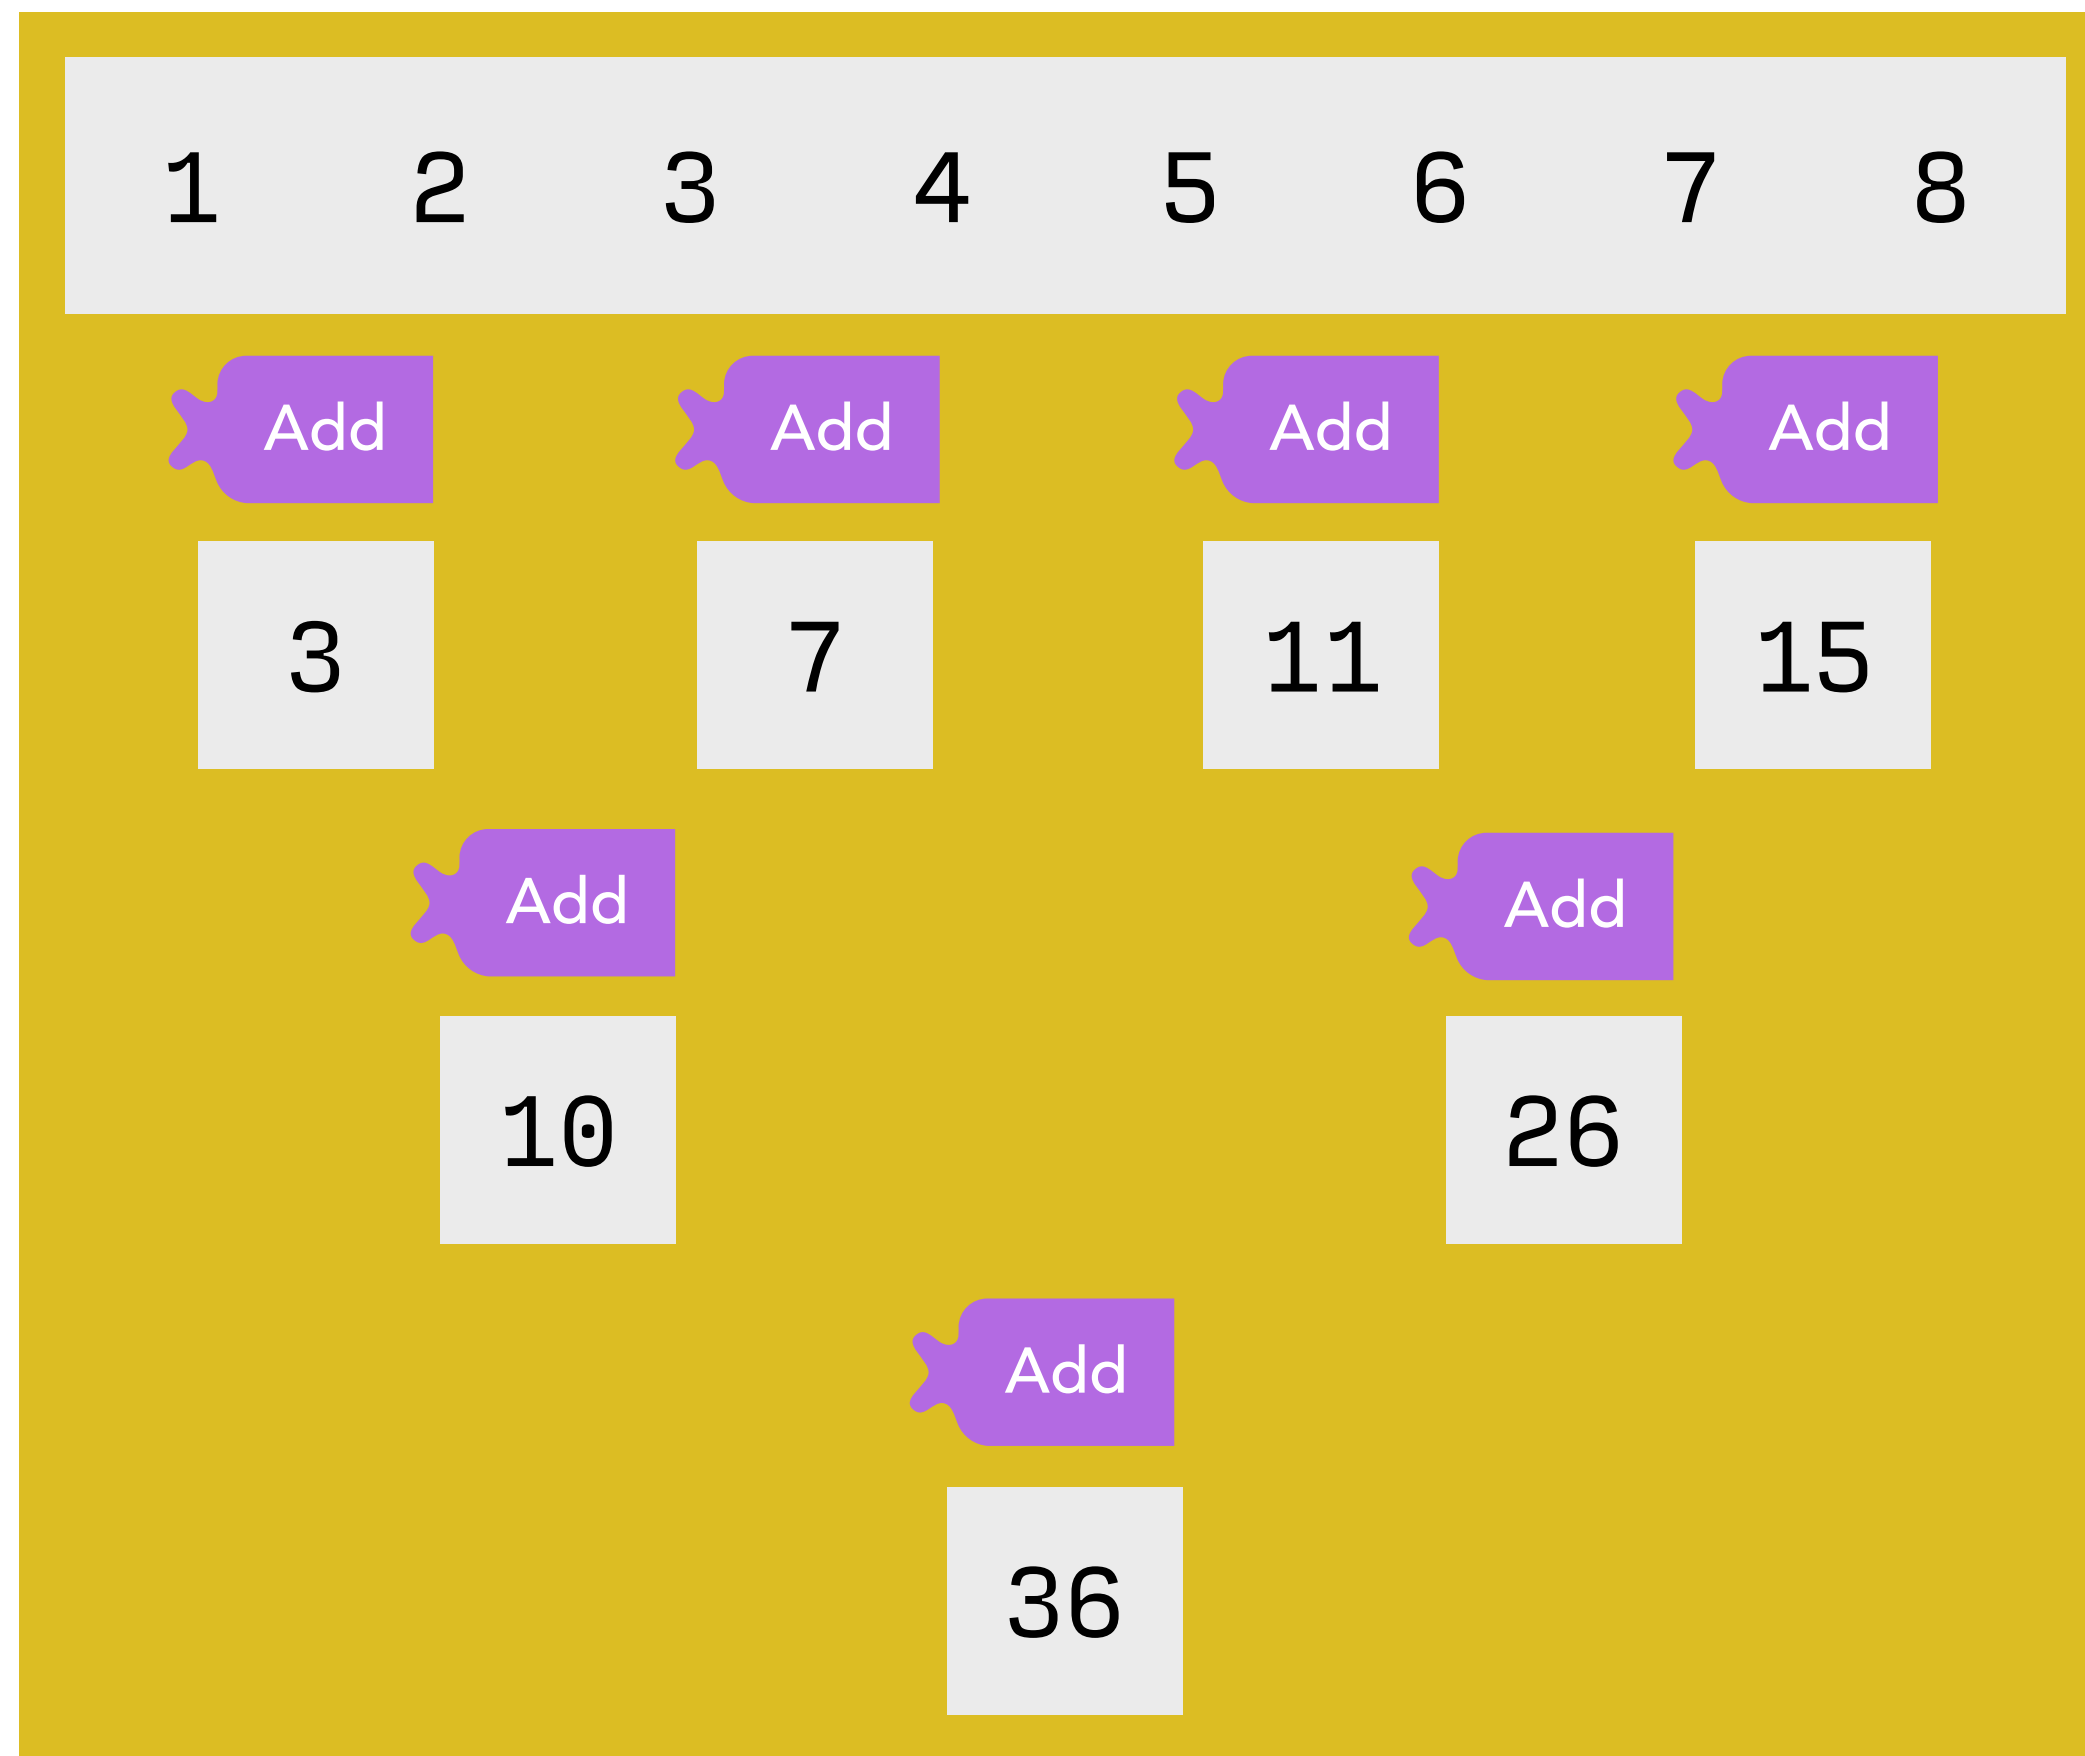
Parallel algorithm?



Dependencies!



Parallel algorithm (assuming associativity)



SkePU Introduction and History

- SkePU uses "modern" C++

```
// "auto" type specifier
auto addOneMap = skepu::Map<1>(addOneFunc);

skepu::Vector<float> input(size), res(size);
input.randomize(0, 9);

// Lambda expression
auto dur = skepu::benchmark::measureExecTime([&]
{
    addOneMap(res, input);
});
```

capture by
reference



- Implementation reliant on variadic templates, template metaprogramming, and other C++11 features

- ✓ Efficient parallel algorithms
- ✓ Accessible interface
- ✓ Memory management and data movement
- ✓ Automatic backend selection and tuning

- **SkePU 1**, public release **2010**

Enmyren, Kessler, HLPP 2010
<https://doi.org/10.1145/1863482.1863487>
 (See website for more publications)

- C++ template-based interface (limited arity)
- Multi-backend using macro-based code generation

- **SkePU 2**, public release **2016**

Ernstsson, Li, Kessler, Int J Parallel Prog 46, 62–80 (2018)
<https://doi.org/10.1007/s10766-017-0490-5>

- C++11 *variadic* template interface (flexible arity)
- Multi-backend using source-to-source precompiler

- **SkePU 2.1 (2017)** *Experimental feature*: Lazy evaluation

Ernstsson, Kessler,
 Concurrency Computat Pract Exper. 2019; 31:e5003
<https://doi.org/10.1002/cpe.5003>

- **SkePU 2.2**, public release **2018**

- Hybrid CPU-GPU backends

Öhberg, Ernstsson, Kessler,
 J Supercomput 76, 5038–5056 (2020)
<https://doi.org/10.1007/s11227-019-02824-7>

- **SkePU 2.3 (2019)** *Experimental feature*: Multi-variant user functions

Ernstsson, Kessler,
 Parallel Computing: Technology Trends,
 IOS PRESS , 2020, Vol. 36, pp. 475-484
<http://doi.org/10.3233/APC200074>

- **SkePU 3**, public release **2020**: Expanding skeleton set, container set, and expressivity

- MapPairs, MapPairsReduce
- Tensor containers (3D and 4D) and new “proxy” containers MatRow, MatCol
- Cluster backend with StarPU-MPI
- Improved syntax and memory consistency model
- Dynamic scheduling

Ernstsson, Ahlqvist, Zouzoula, Kessler,
Int J Parallel Prog (2021).
<https://doi.org/10.1007/s10766-021-00704-3>

Panagiotou, Ernstsson, Ahlqvist, Papadopoulos, Kessler, Soudris,
SCOPES '20 proceedings, Pages 74–77
<https://doi.org/10.1145/3378678.3391889>

Papadopoulos *et al.*,
IEEE Transactions on Parallel and Distributed Systems
<https://doi.org/10.1109/TPDS.2021.3104257>

- **SkePU 3.1**, public release **2021**: SkePU “standard library”

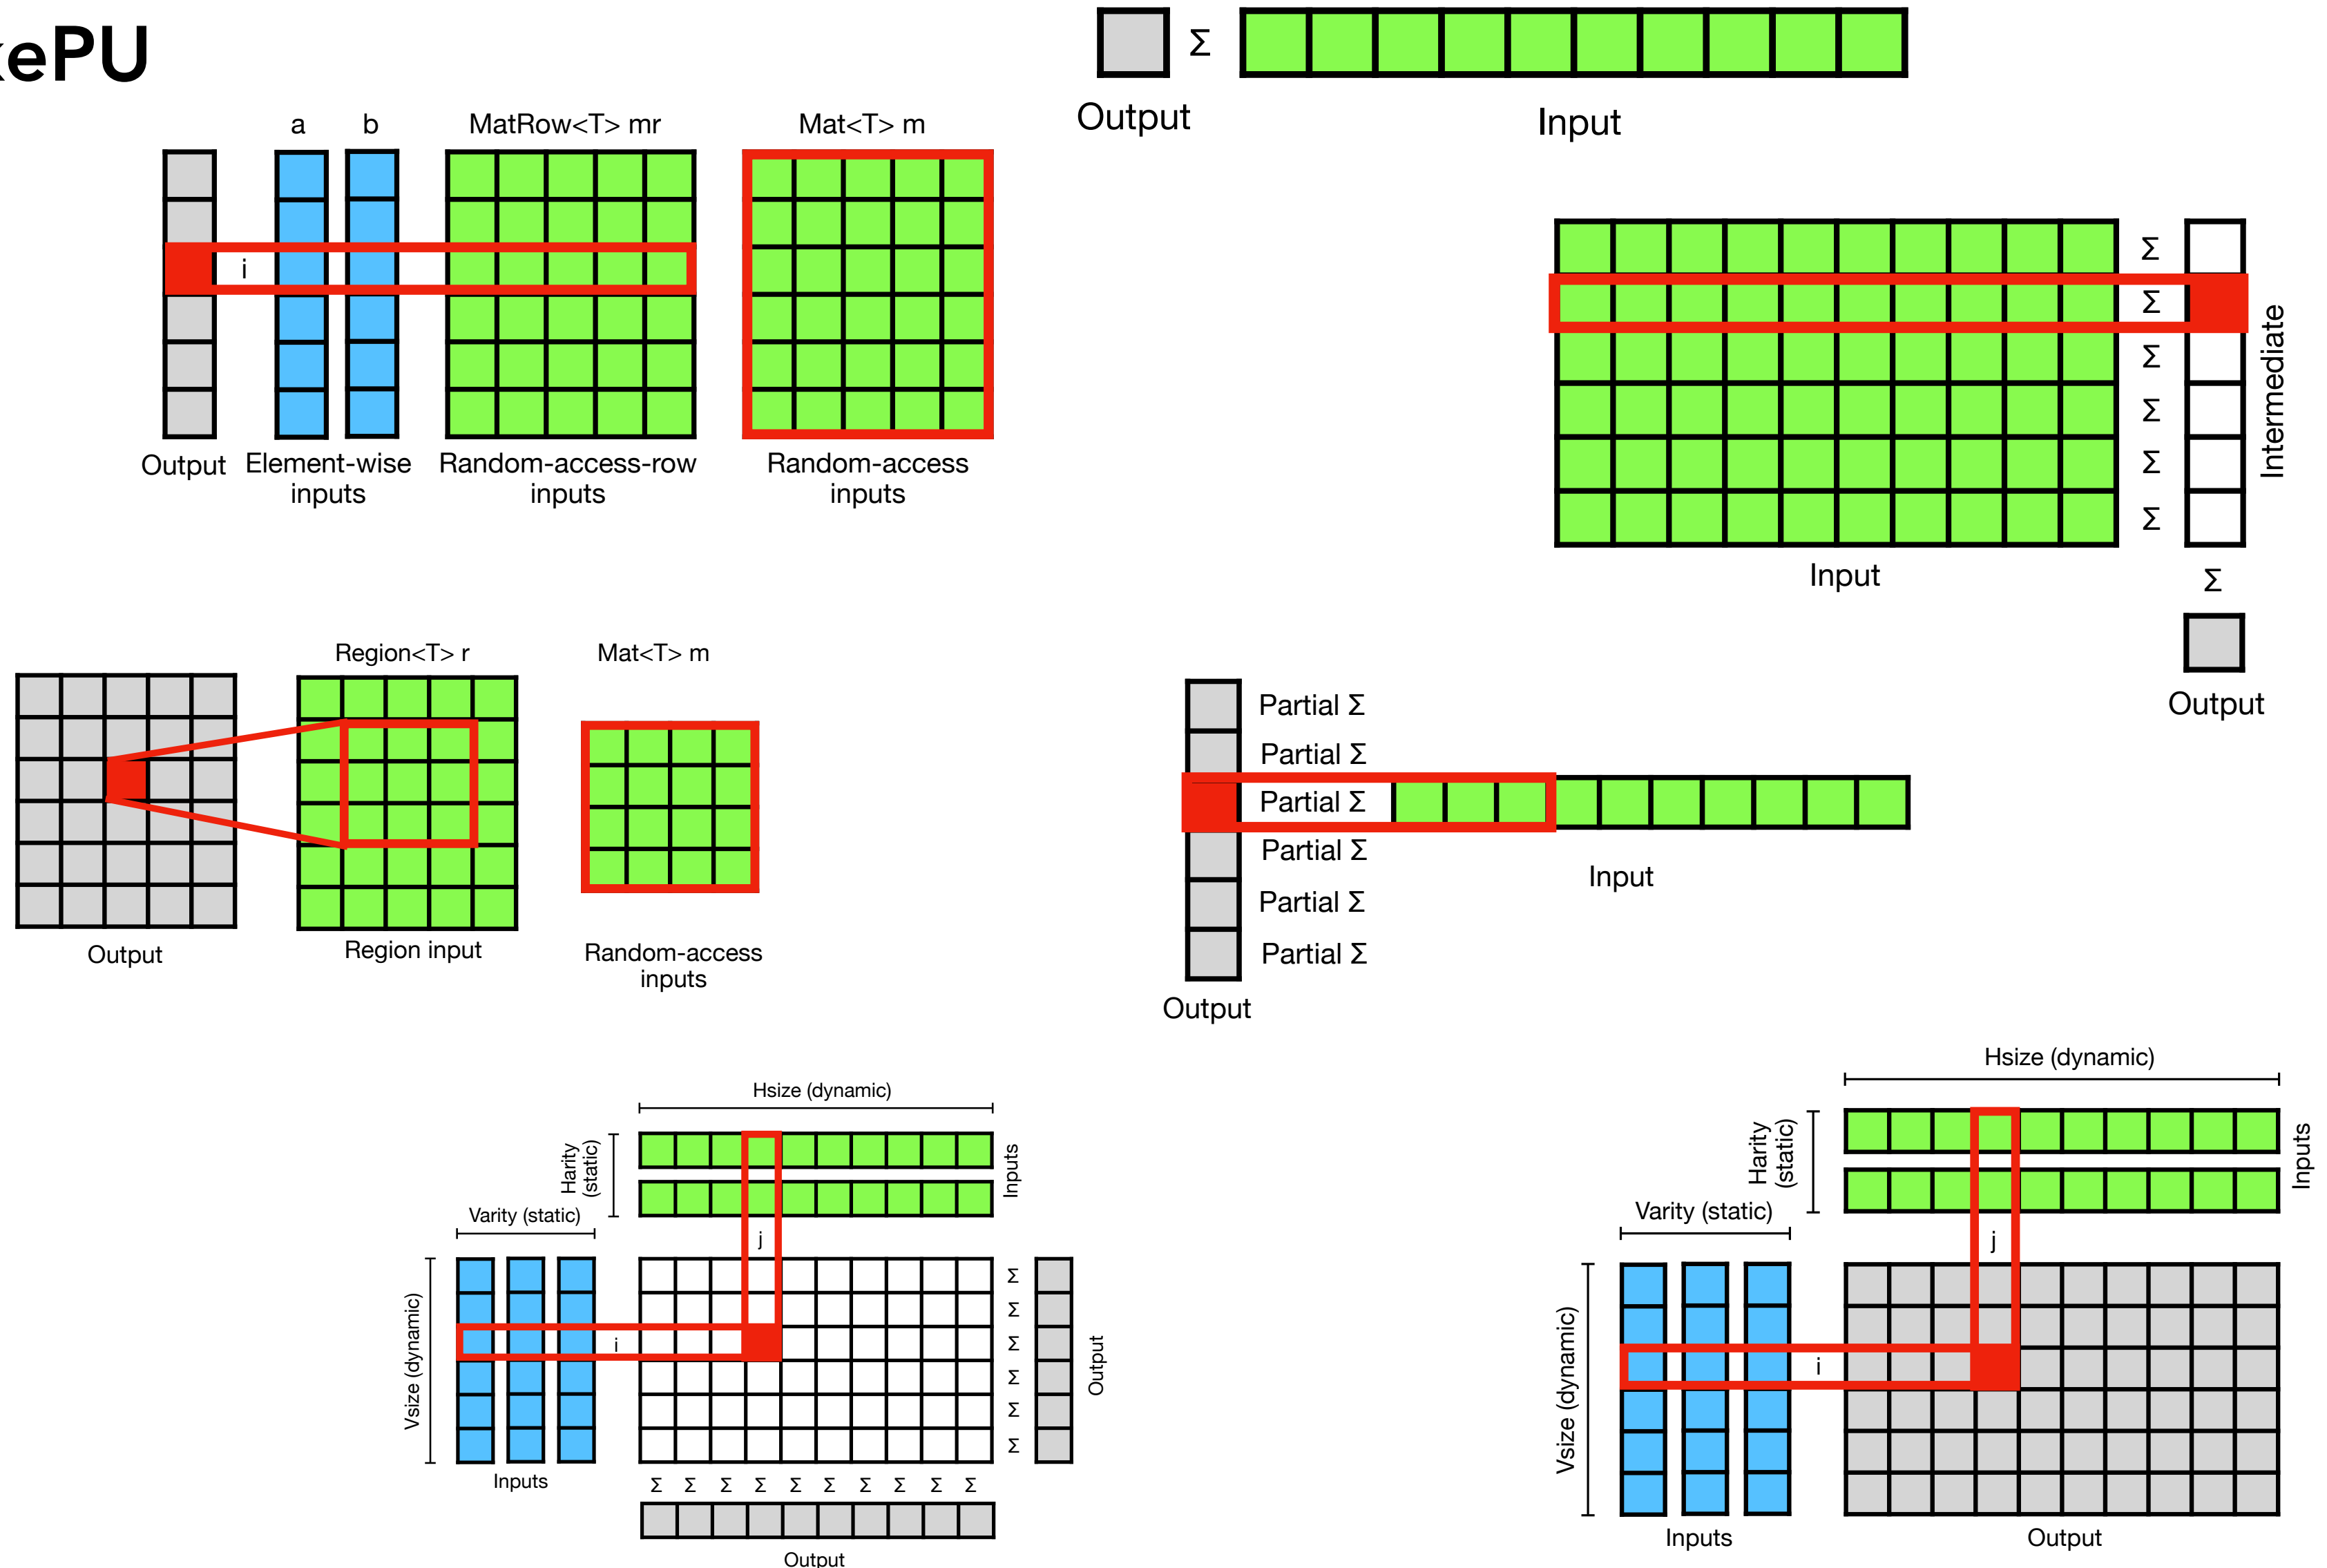
- Complex number API
- BLAS (level 1 + dense level 2, 3)
- Deterministic PRNG
- Strided Map skeletons

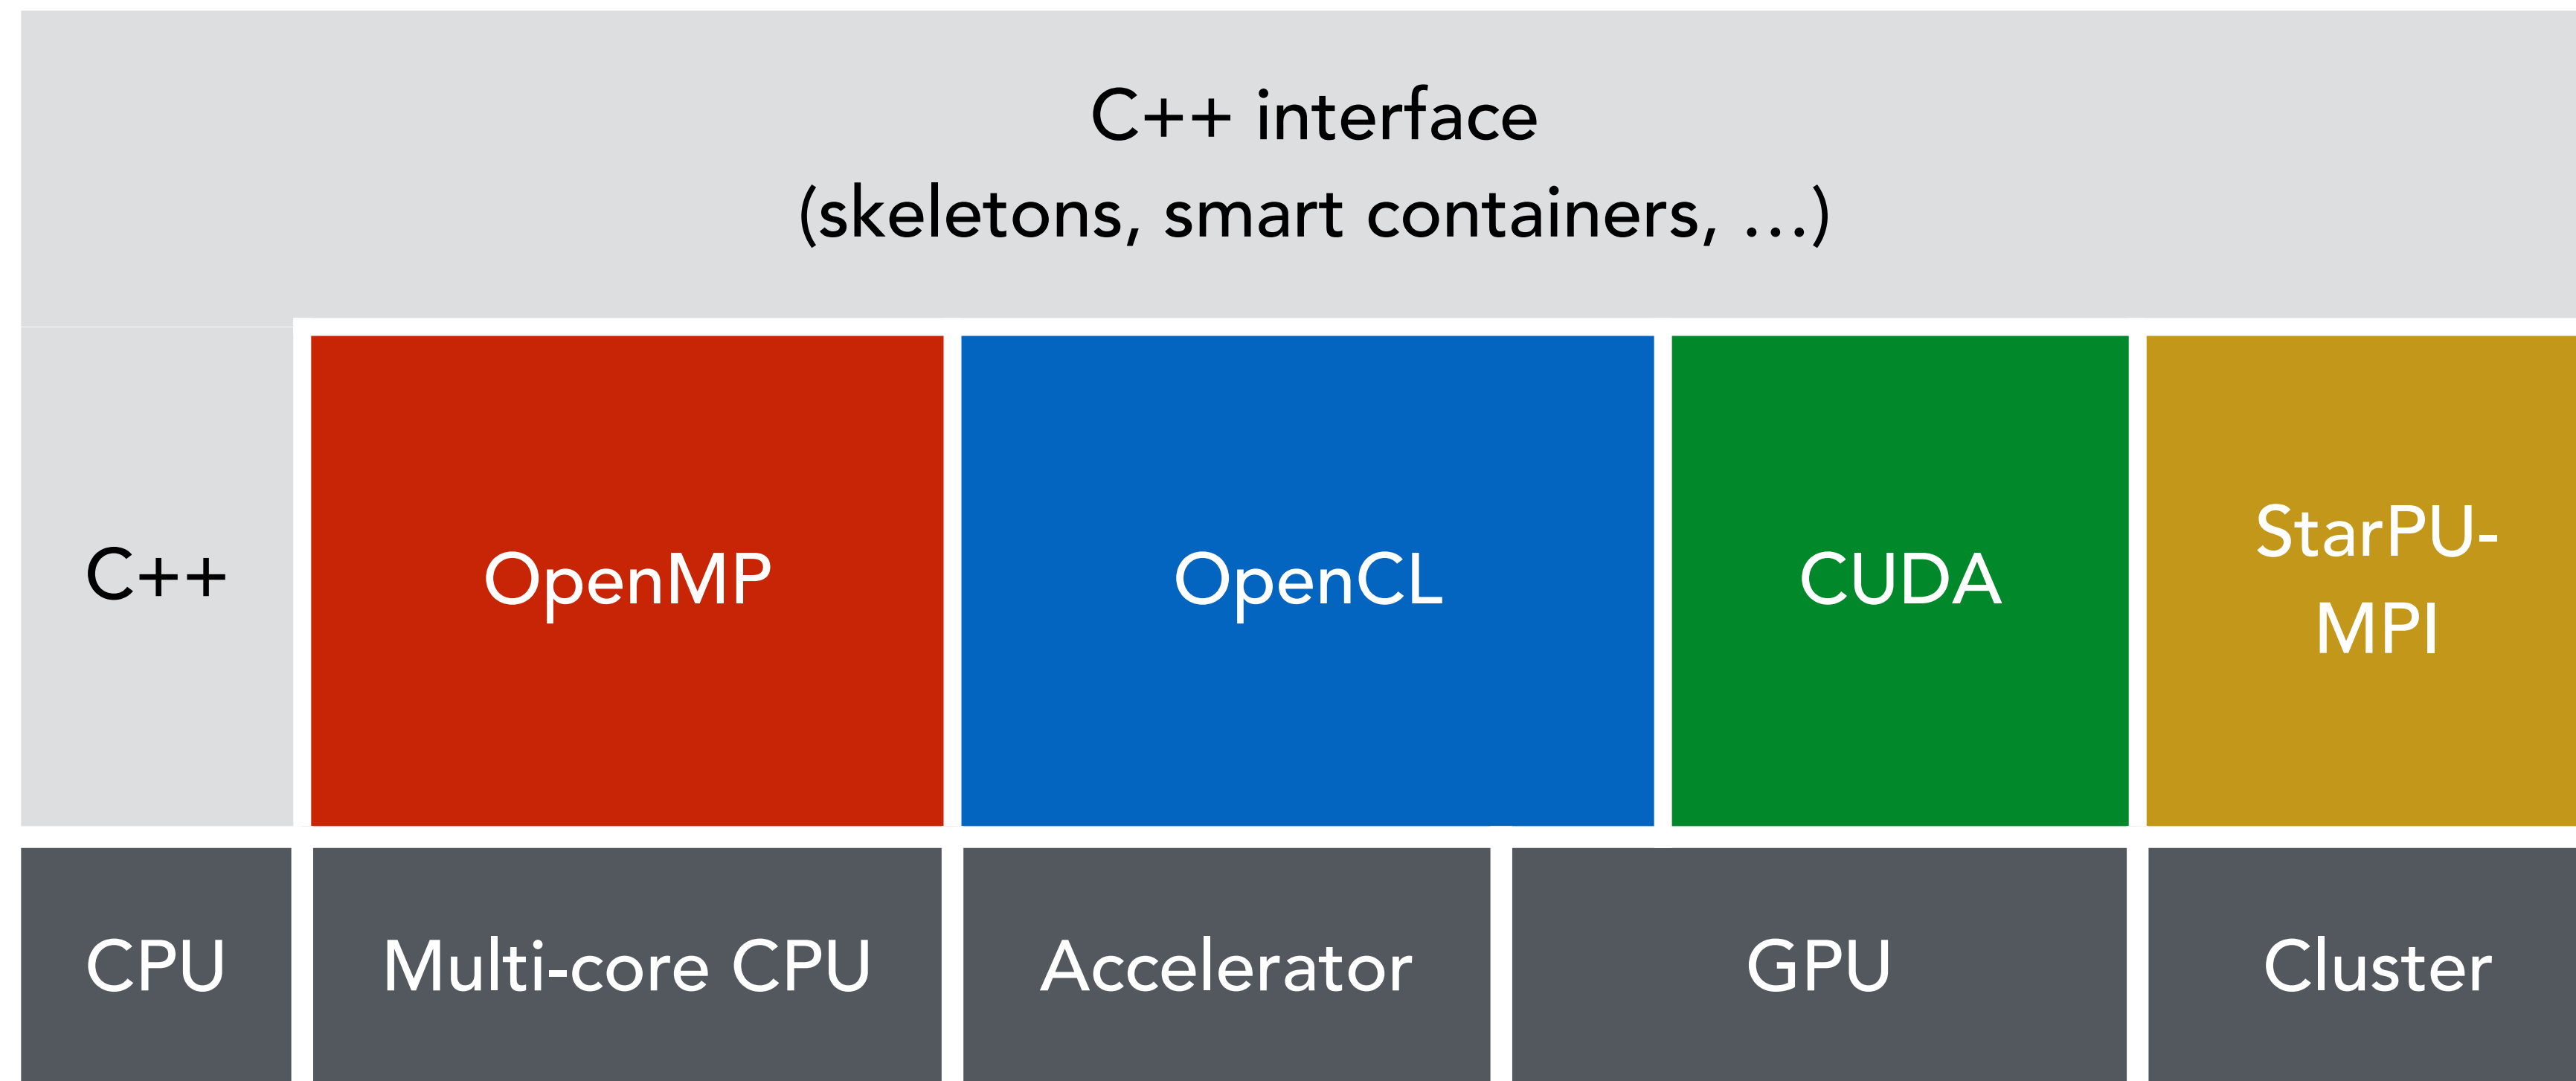
Ernstsson, Kessler, HLPP 2021, to appear.

- Skeleton programming framework
 - C++11 **library** with skeleton and data container classes
 - A **Clang**-based source-to-source **pre-compiler**
- Smart containers: `Vector<T>`, `Matrix<T>`, `Tensor3<T>`, `Tensor4<T>`
 - In development: `SparseMatrix<T>`
- For **heterogeneous multicore** systems
 - Multiple backends
- Active research tool with a number of publications 2010-2022 (see website)

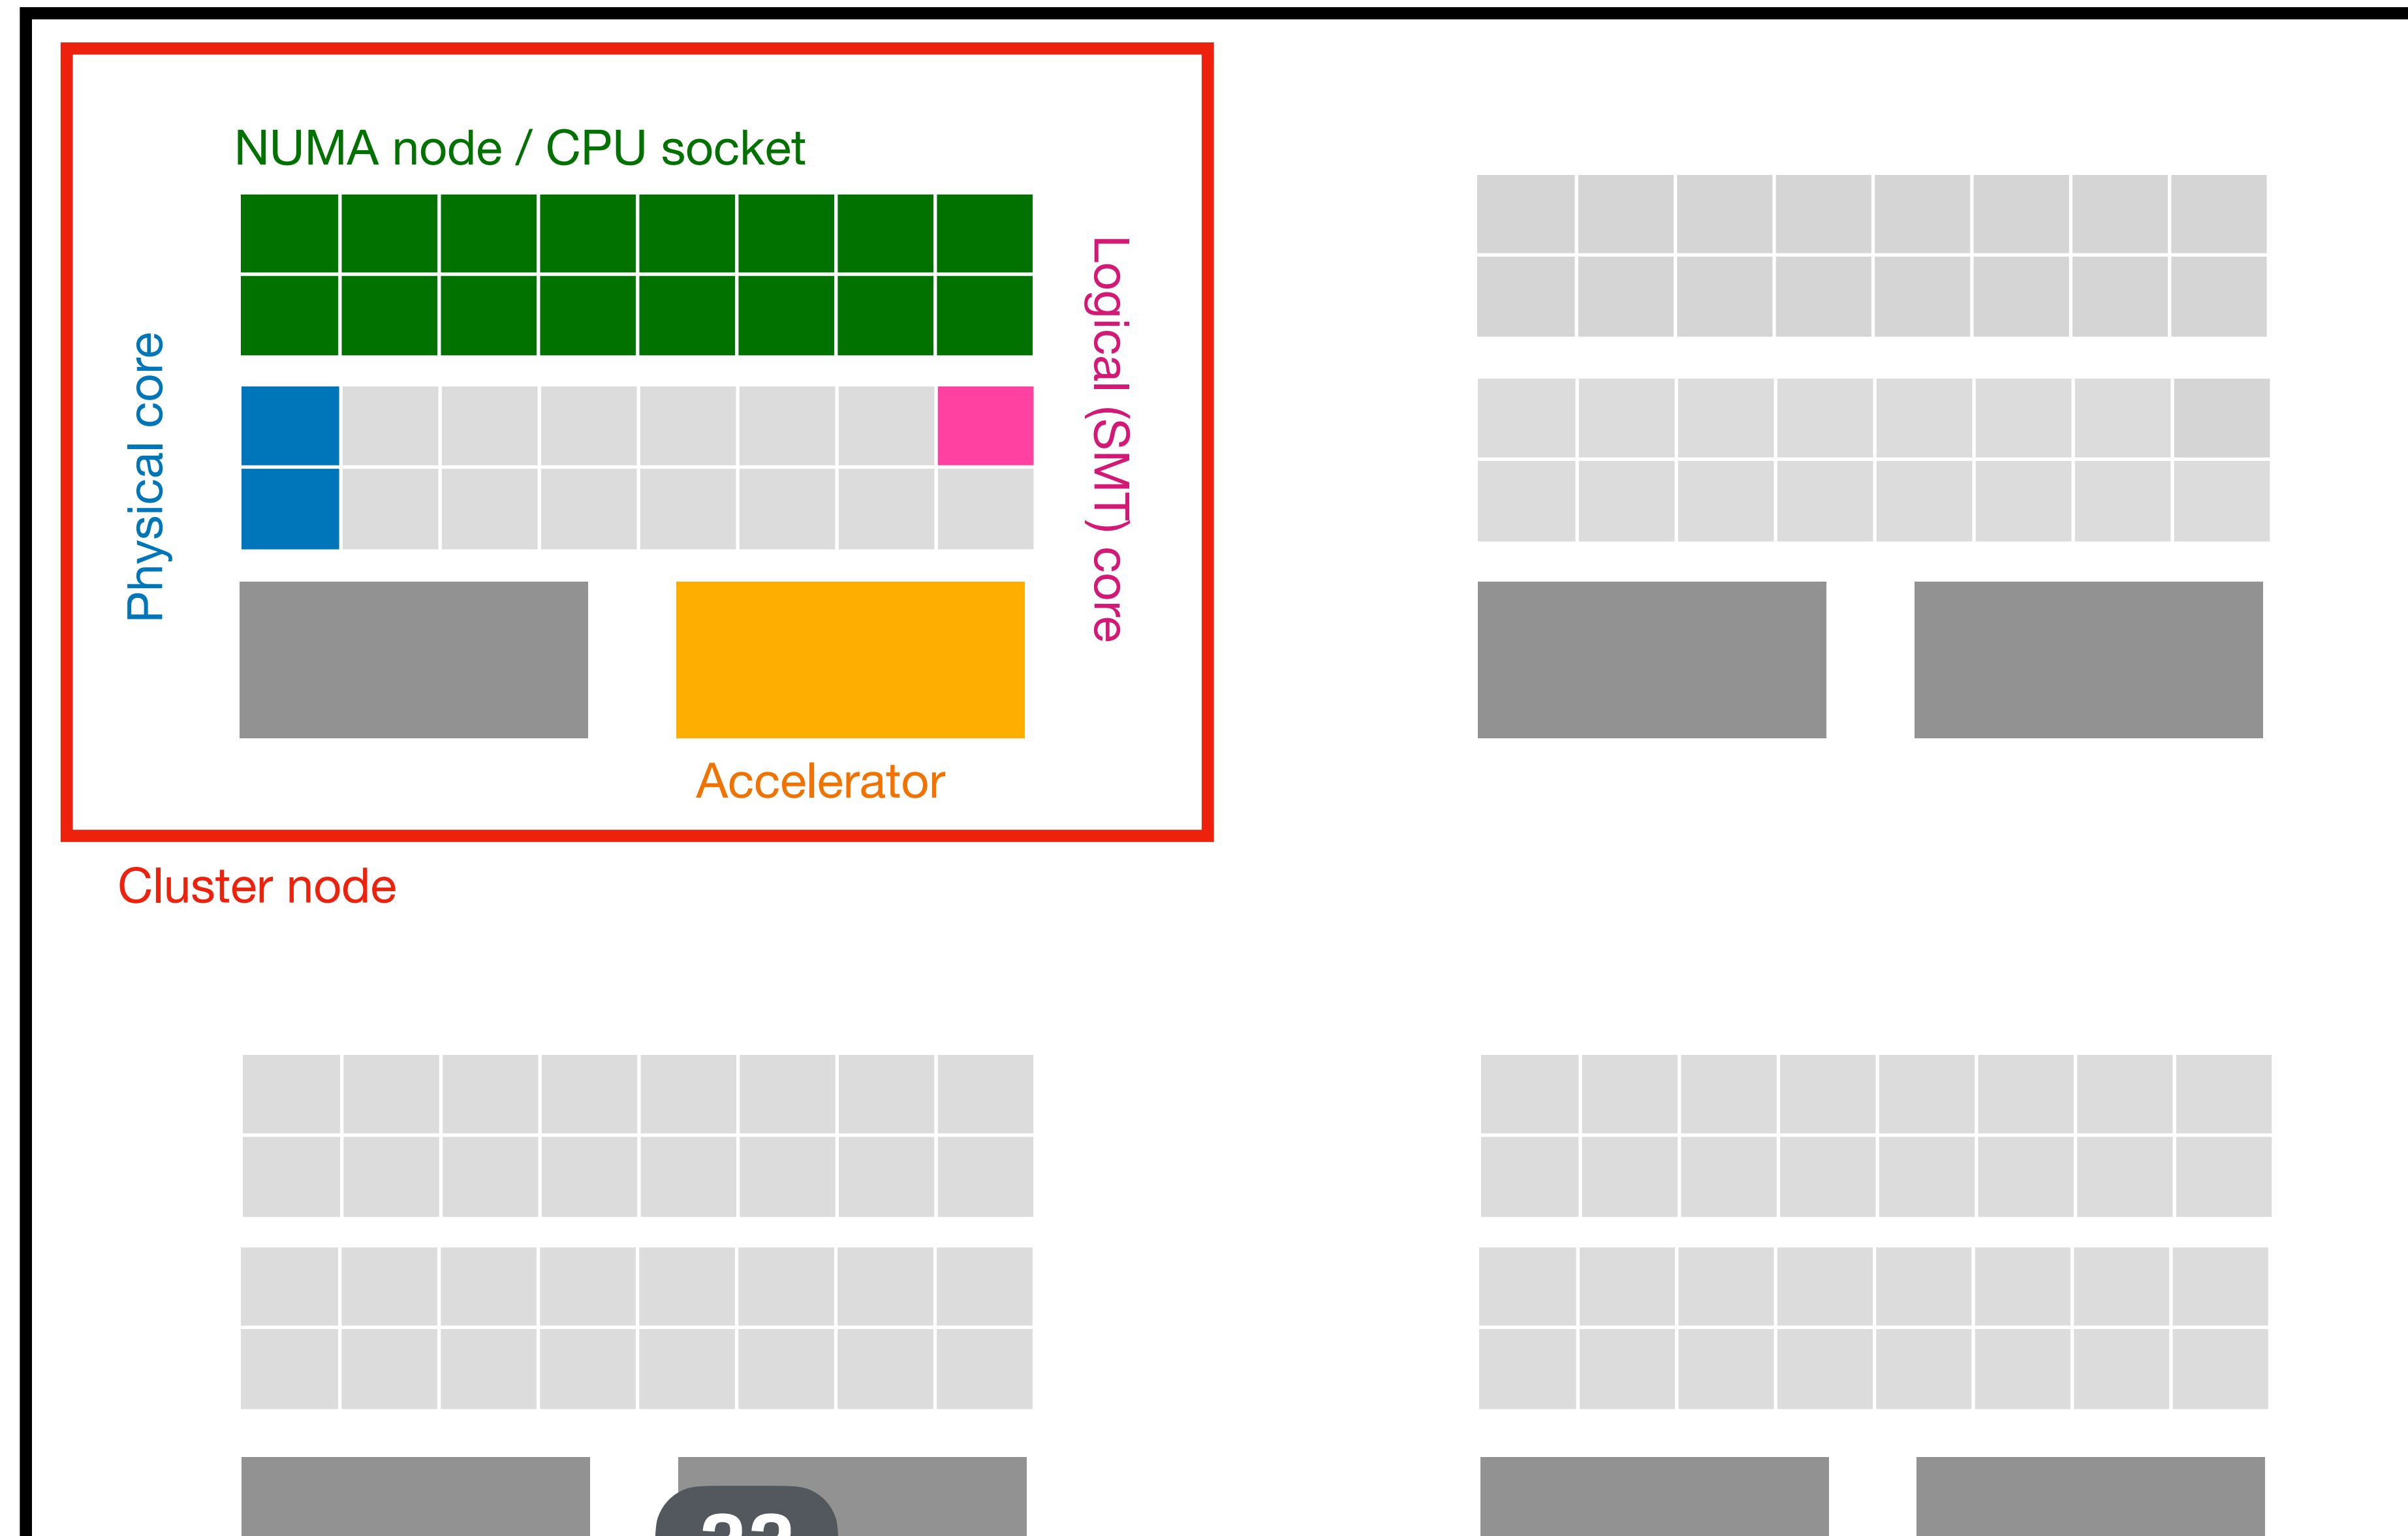
• Skeletons provided by SkePU

- Map
- Reduce
- MapReduce
- Scan
- MapOverlap
- MapPairs
- MapPairsReduce





Cluster




Compute


Management

Idle


Backend **CPU**




Backend **OpenMP**



Backend **OpenCL**
Backend **CUDA***



Backend **Hybrid**




*) NVIDIA GPU accelerators only

Compute


Management

Idle

Backend **StarPU-MPI**



Backend **StarPU-MPI-CUDA***



*) NVIDIA GPU accelerators only

Vector

i	0	1	2	3	4
	0	1	2	3	4

Matrix

j	0	1	2	3	4
i					
0	0	1	2	3	4
1	5	6	7	8	9
2	10	11	12	13	14
3	15	16	17	18	19
4	20	21	22	23	24


i	0			1			
j	k	0	1	2	0	1	2
0		0	1	2	9	10	11
1		3	4	5	12	13	14
2		6	7	8	15	16	17

Tensor3


		j						
		0			1			
		1	0	1	2	0	1	2
i	k							
0	0	0	1	2	9	10	11	
	1	3	4	5	12	13	14	
	2	6	7	8	15	16	17	
1	0	18	19	20	27	28	29	
	1	21	22	23	30	31	32	
	2	24	25	26	33	34	35	

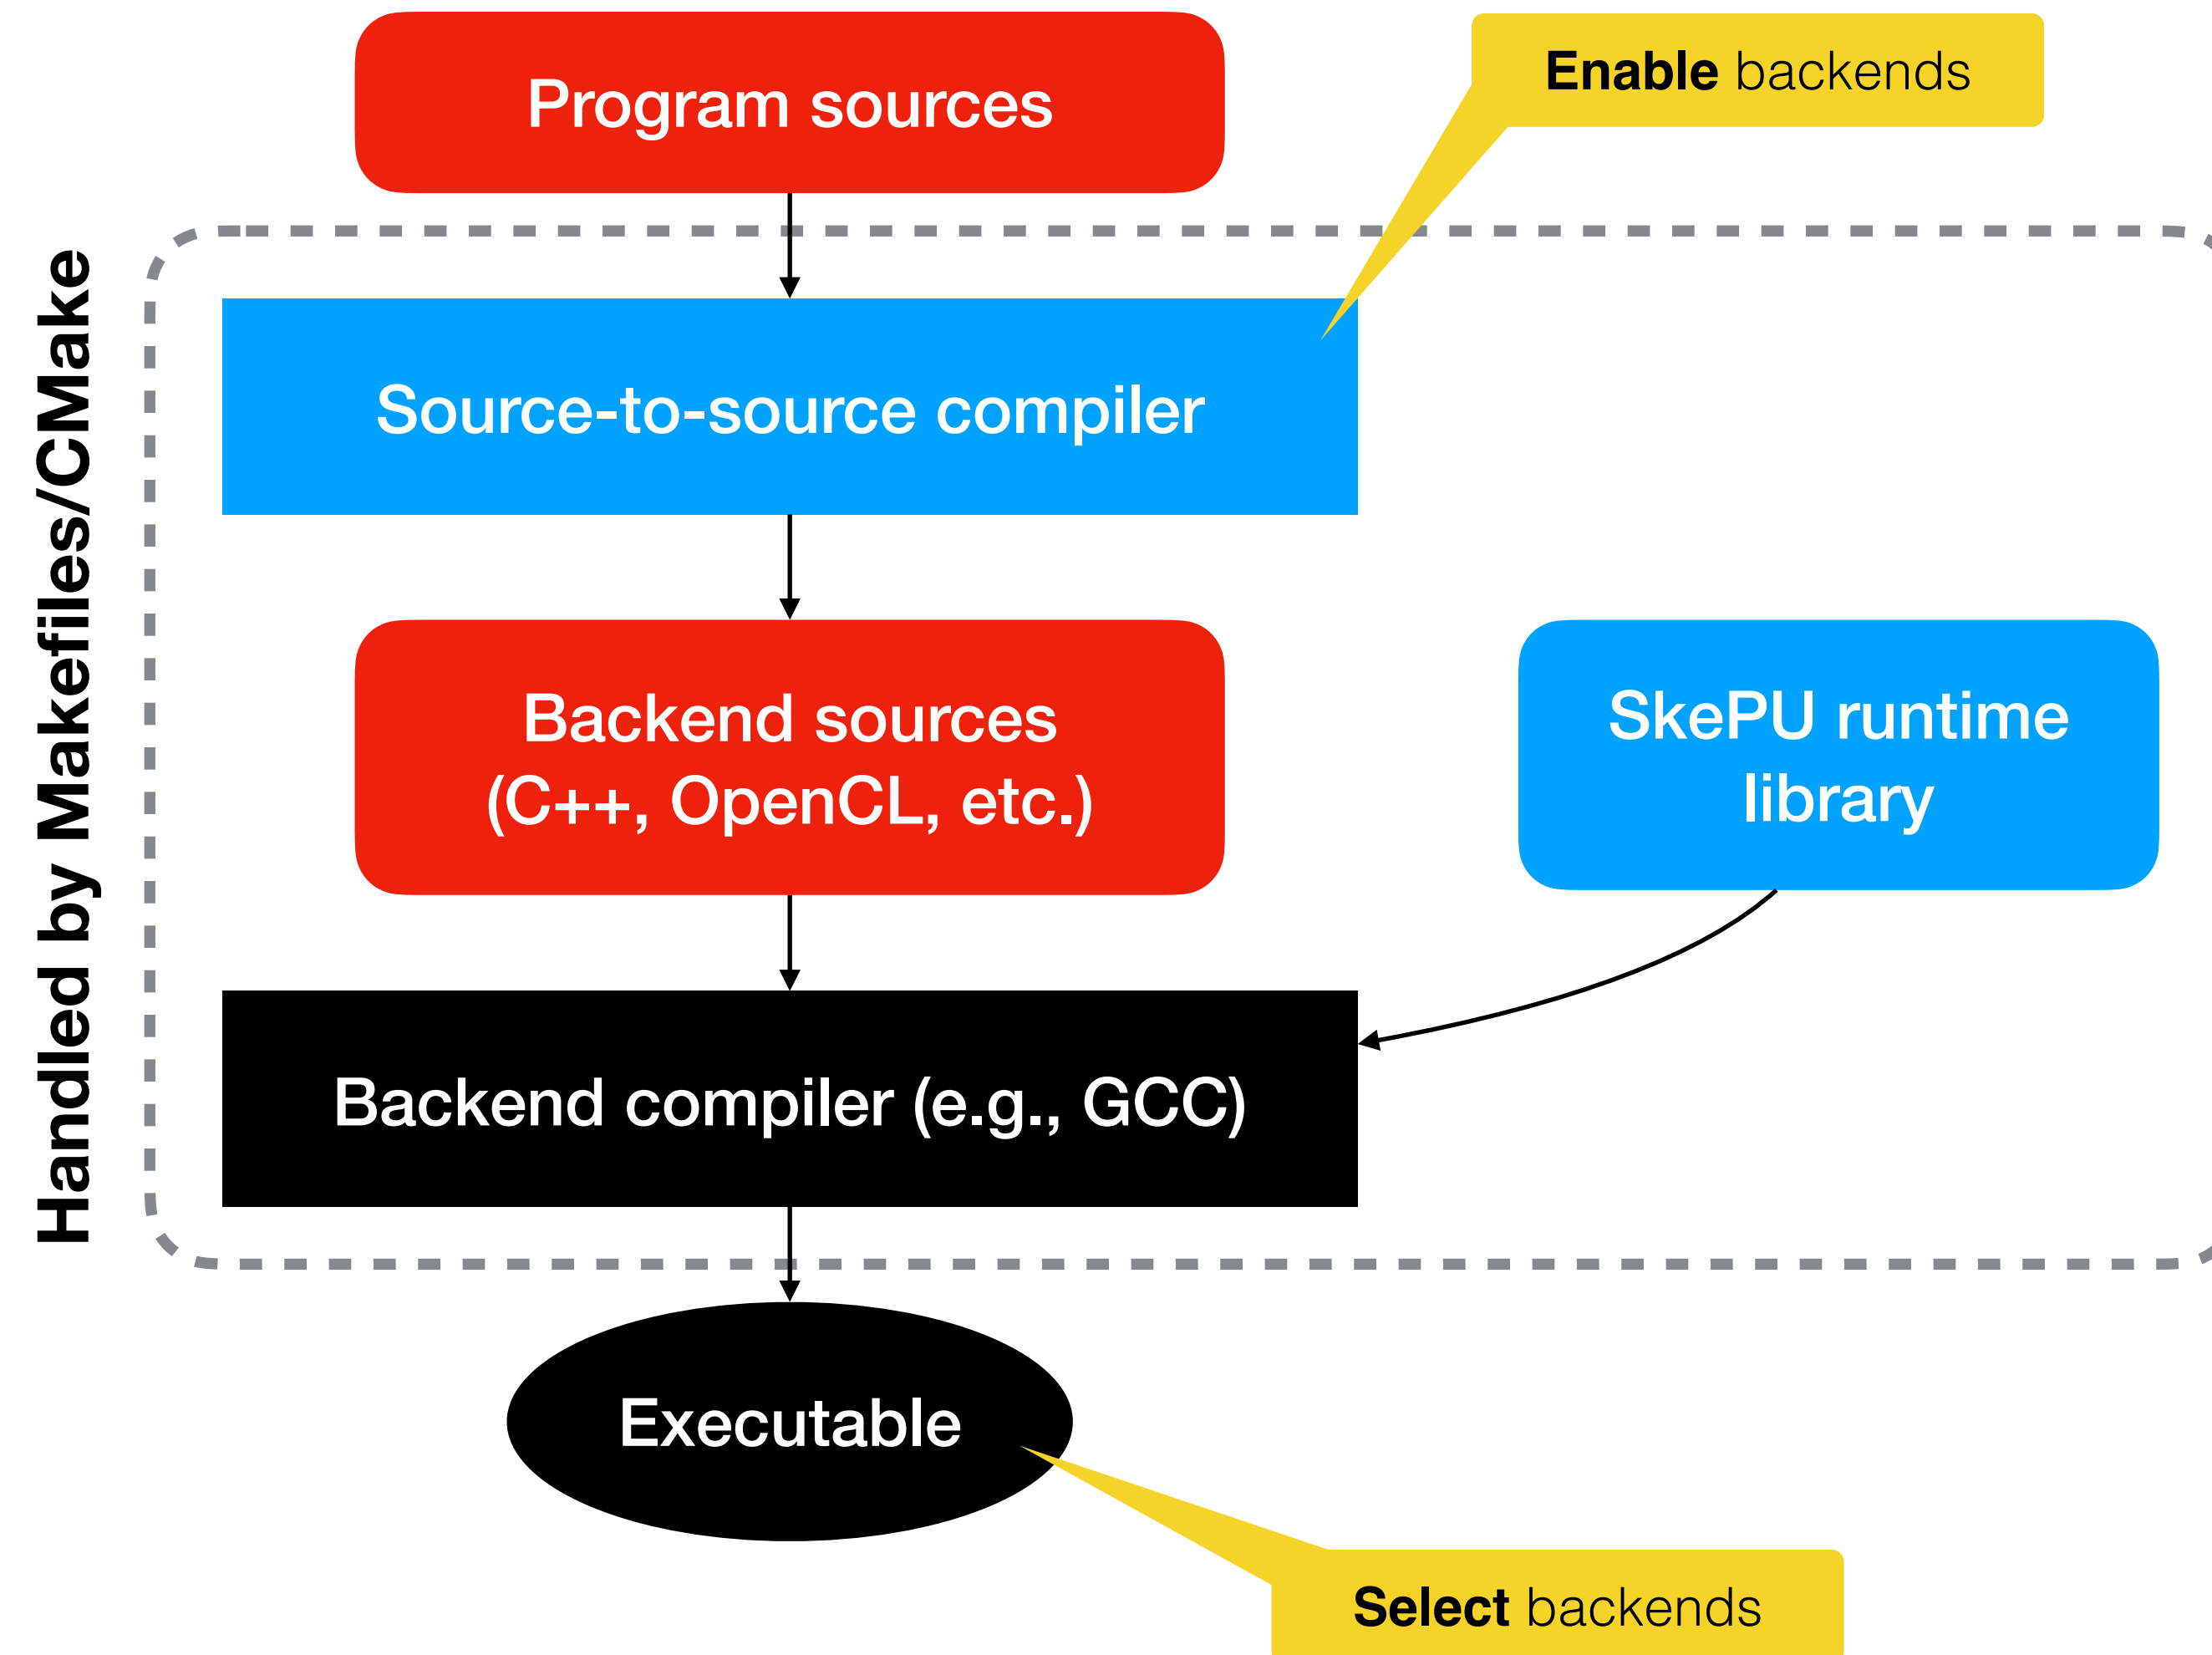
Tensor4

- C++ template class instance
- Contains:
 - CPU memory buffer **pointer** (alt. StarPU handles)
 - Size information (size, width/height)
 - OpenCL/CUDA/MPI **handles**
 - Consistency states
- Template type can be **custom struct**, but be careful!
 - Data layout not verified across backends/languages



Using SkePU







```
int add(int a, int b, int m)
{
    return (a + b) % m;
}
```



```
auto vec_sum = Map<2>(add);
```



```
vec_sum(result, v1, v2, 5);
```



- User functions are C++ (rather, C) functions
- The signature is analyzed by the pre-compiler to extract the skeleton signature
- Each skeleton has their own expected patterns for UF parameters (but the general structure is shared)
- The UF body is **side-effect free** C (compatible with CUDA/OpenCL)
 - No communication/synchronization
 - No memory allocation
 - No disk IO

```
int add(int a, int b, int m)
{
    return (a + b) % m;
}
```

- **Variable arity** on Map and MapReduce skeletons
- **Index** argument (of current Map'd container element)
- **Uniform** arguments
- Smart container arguments accessible **freely** inside user function
 - **Read**-only / **write**-only / **read-write** copy modes
- User function **templates**


```
template<typename T>
T abs(T input)
{
    return input < 0 ? -input : input;
}
```

```
template<typename T>
T user_function(Index1D row, const Mat<T> m, const Vec<T> v)
{
    T res = 0;
    for (size_t i = 0; i < v.size; ++i)
        res += m(row.i * m.cols + i) * v(i);

    return abs(res);
}
```

- Multi-variant user function **specialization**
- Targeting backend
- Custom **types**
- **Chained** user functions
- In-line **lambda** syntax for user functions
- **"Intrinsic"** functions
Some functions exist in the standard library of all SkePU backends.
Examples: $\sin(x)$, $\text{pow}(x, e)$

```
auto vec_sum = skepu::Map([](int a, int b)
{
    return a + b;
});

// ...

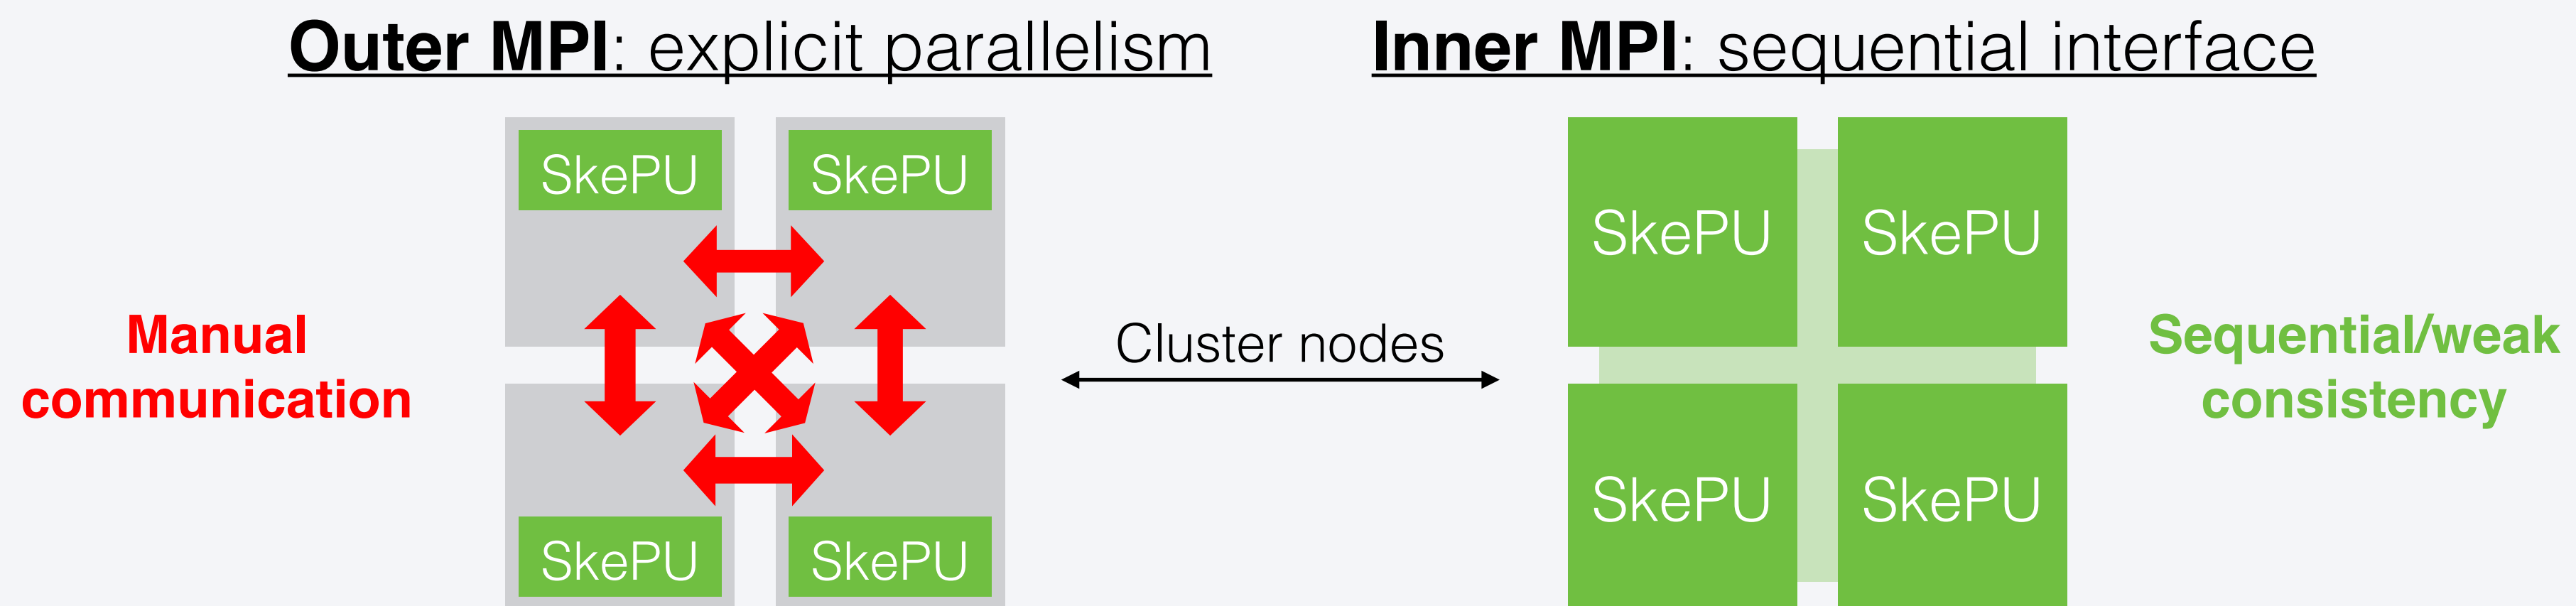
vec_sum(result, v1, v2);
```

Compilation Options and Backend Selection

- Precompiler options
 - `skepu-tool`
`map.cpp`
`-name map_precompiled`
`-dir bin`
`-openmp -opencl`
`-- [Clang flags]`
- Handled by Makefiles in the binary tutorial distribution
- Handled by CMake in standard SkePU repository
- Clang flags:
 - Include path to Clang language-headers
 - Include path to SkePU headers
 - Optional OS-specific paths and settings


- `auto spec = skepu::BackendSpec{argv[2]};`
- Sets the **run-time backend** from a string, must match any of:
 - "CPU", "OpenMP", "OpenCL", "CUDA", "Hybrid"
- Cluster backend is a **compile-time** selection due to SPMD model
 - Uses OpenMP on node level by default
 - Optional GPU backend selection
- Global backend spec for all skeletons, overridden by instance-specific spec.
- `skepu::setGlobalBackendSpec(spec);`

- SkePU can **target clusters**
 - Any valid SkePU 3 program needs not be syntactically changed to use the cluster backend ("inner MPI mode")
 - SkePU can also run on clusters by using skeletons locally on each node ("outer MPI mode": MPI+SkePU)
- Implementation atop *StarPU*
 - Challenges: SPMD-style; C-level APIs (StarPU + MPI)



- **Backend specifications** can furthermore contain execution parameters
 - Number of **threads**
 - Number of **GPUs**
 - GPU block size, thread count hinting
 - OpenMP **scheduling** mode
 - Static (default) / dynamic / guided / auto
 - Chunk size for dynamic scheduling
 - **Partition ratio** for Hybrid backend

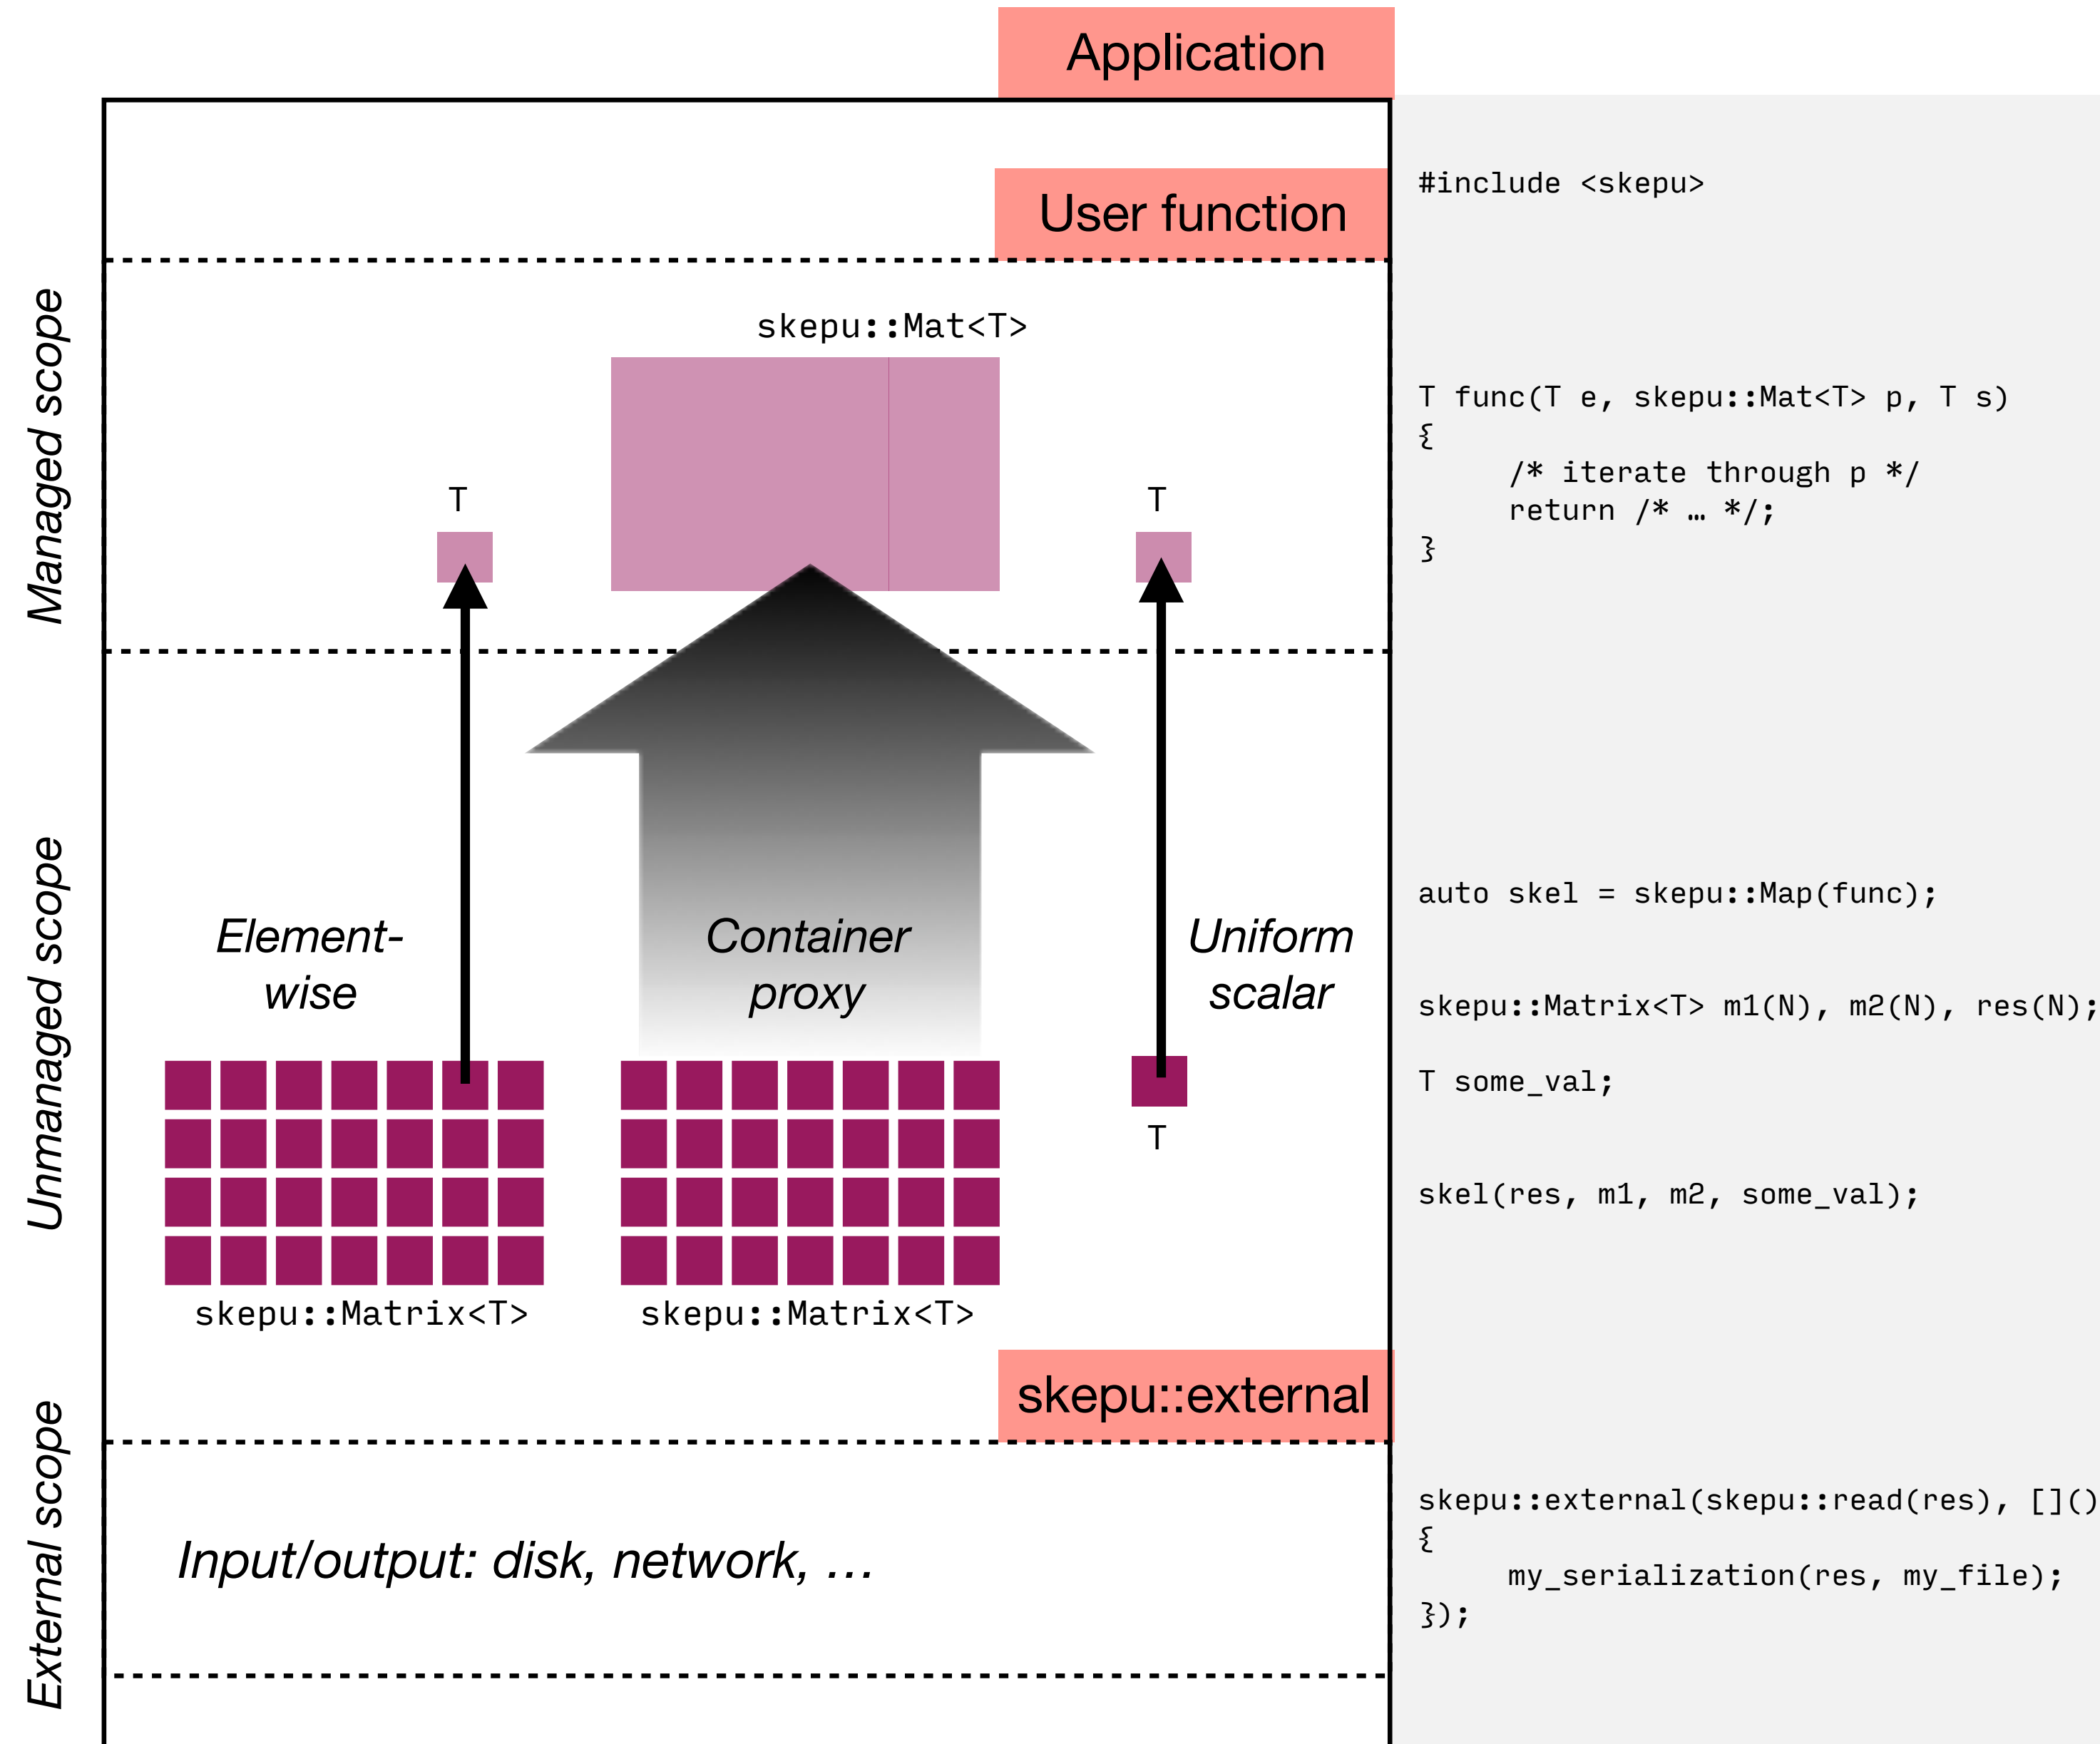
- Three benchmarks that illustrate the effects of dynamic scheduling in OpenMP backend
- All variants are SkePU backend executions
- Sequential CPU and OpenCL included for reference



Smart Containers & Consistency Model

- SkePU provides three different "scopes" with different syntax and consistency rules
- Inside a user function: **Managed scope**
 - C-style syntax allowed, no global synchronization (compare with GPU kernel code)
- Outside user function: **Unmanaged scope**
 - Regular single-threaded C++-land with smart container objects providing weak consistency
- For global side effects: **External scope**
 - In SPMD model (clusters), the C++ program does not run single-threaded
 - External construct wraps file read/write, logging, etc.

Consistency Model



- `container.flush();`
 - Flushes and ensures that the CPU buffer contains up-to-date values.
- `container(i) = value;`
 - Direct element access into the raw CPU buffer.
 - Weak consistency, unless compile-time flag enabled (optional for debugging purposes)
- `skepu::external(skepu::read(particles), [&]{
 save_step(particles, outfile);
});`
 - Manages IO and other external operations in SPMD mode

SkePU **DEMO**

Skeletons In Depth

Feature \ Skeleton	Map	MapPairs	MapOverlap	Reduce	Scan	MapReduce	MapPairsReduce
Elwise dimension in	1–4	1	1–4	1–4**	1	1–2	1
Elwise dimension out	Same as <i>in</i>	2	Same as <i>in</i>	0–1	1	0	1
Indexed	Yes	Yes	Yes	-	-	Yes	Yes
Multi-return	Yes	Yes	Yes	-	-	Yes	Yes
Elwise parameters	Variadic	Variadic x2	1***	*	*	Variadic	Variadic x2
Full proxy parameters	Variadic	Variadic	Variadic	-	-	Variadic	Variadic
Uniform parameters	Variadic	Variadic	Variadic	-	-	Variadic	Variadic
Region proxy	-	-	Yes	-	-	-	-
MatRow/MatCol proxy	Yes	Yes	-	-	-	Yes	Yes
Footnotes							
*	Parameters to the user functions can be raw elements from the container or partial results, depending on evaluation order.						
**	Dimensions higher than 2 are linearized in the current implementation.						
***	A region of elements surrounding the current index is supplied.						

Map

- Three groups of user function parameters:

- **Element-wise**

Only one element per user function call

- **Random-access containers**

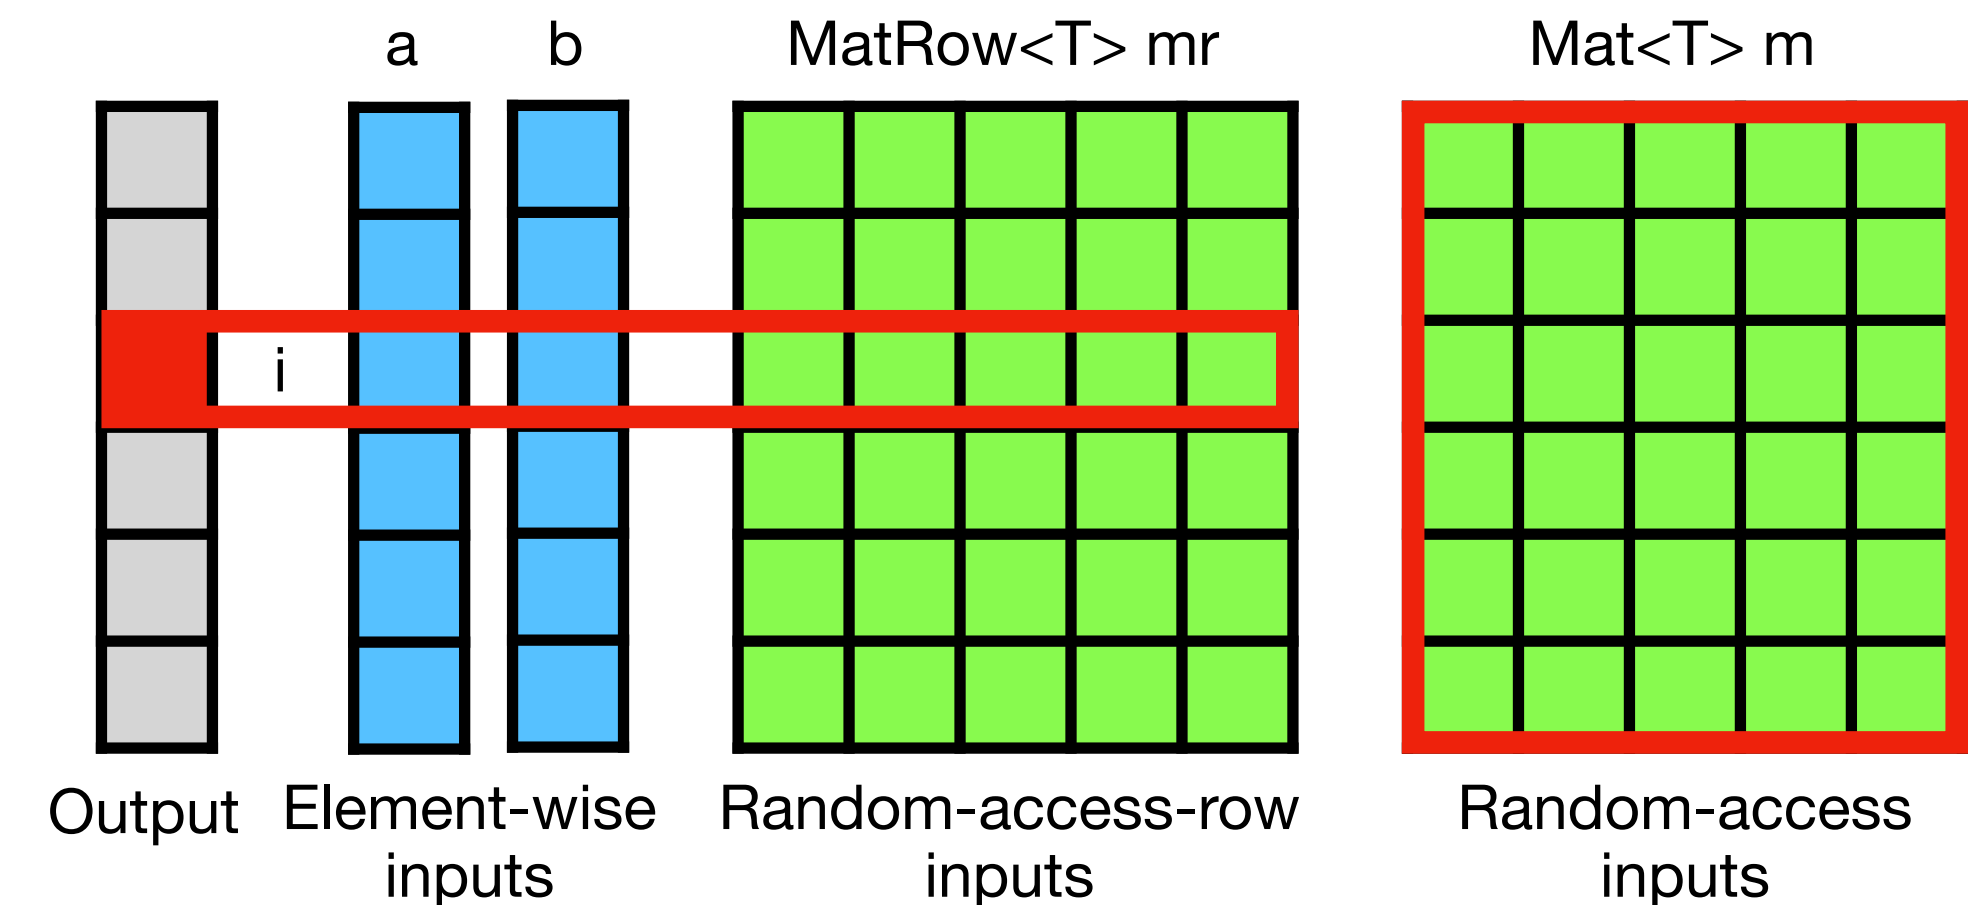
Replicated for each memory space (e.g. GPUs)

Proxy types `Vec<T>` and `Mat<T>` in user function

- **Uniform scalars**

Same values everywhere

- Argument groups are variadic (flexible count, including 0)
- Above order must be obeyed (element-wise first etc.)
- The parallelism/number of user function invocations is always determined by the return container (first argument), also in case of element-wise arity of 0.
- Also applies to **MapReduce**, **MapOverlap**, **MapPairs**, **MapPairsReduce**!



```
float sum(float a, float b)
{
    return a + b;
}
```

```
Vector<float> vector_sum(Vector<float> &v1, Vector<float> &v2)
{
    auto vsum = Map<2>(sum);
    Vector<float> result(v1.size());
    return vsum(result, v1, v2);
}
```

- Map (and MapReduce) can index containers in a **strided fashion**
- `instance.setStride(...);`
- Positive strides indicate length of jump **between elements**
- Negative strides in addition **reverses** indexing order
- Unit, non-unit, and negative strides can be mixed in call

```

auto addr = skepu::Map(add);
auto addrsum = skepu::MapReduce(add, add);

constexpr size_t N{8};
skepu::Vector<int> v1(N*4), v2(N*3), r(N*2), r2(N);

// MAP, POSITIVE STRIDES
addr.setStride(2, 4, 3);
addr(r, v1, v2);

// MAPREDUCE, POSITIVE STRIDES
addrsum.setStride(4, 3);
int res = addrsum(v1, v2);


// MAP, NEGATIVE STRIDES
addr.setStride(2, -4, -3);
addr(rc, v1, v2);

```

- Optionally, use iterators with Map (or MapReduce):
 - `mapper(r.begin(), r.end(), v1.begin(), v2.begin());`

Reduce

- **1D Reduce**




- Regular Vector
- Matrix RowWise (returns Vector)
- Matrix ColWise (returns Vector)

- **"2D" Reduce**

```
instance.setReduceMode(ReduceMode::RowWise) // default
instance.setReduceMode(ReduceMode::ColWise)
```

- Regular Matrix (treated as a vector)
- `instance.setStartValue(value)`
- Set Reduction start value. Defaults to 0-initialized.



```
float min_f(float a, float b)
{
    return (a < b) ? a : b;
}
```

```
float min_element(Vector<float> &v)
{
    auto min_calc = Reduce(min_f);
    return min_calc(v);
}
```



```
float plus_f(float a, float b)
{
    return a + b;
}
```

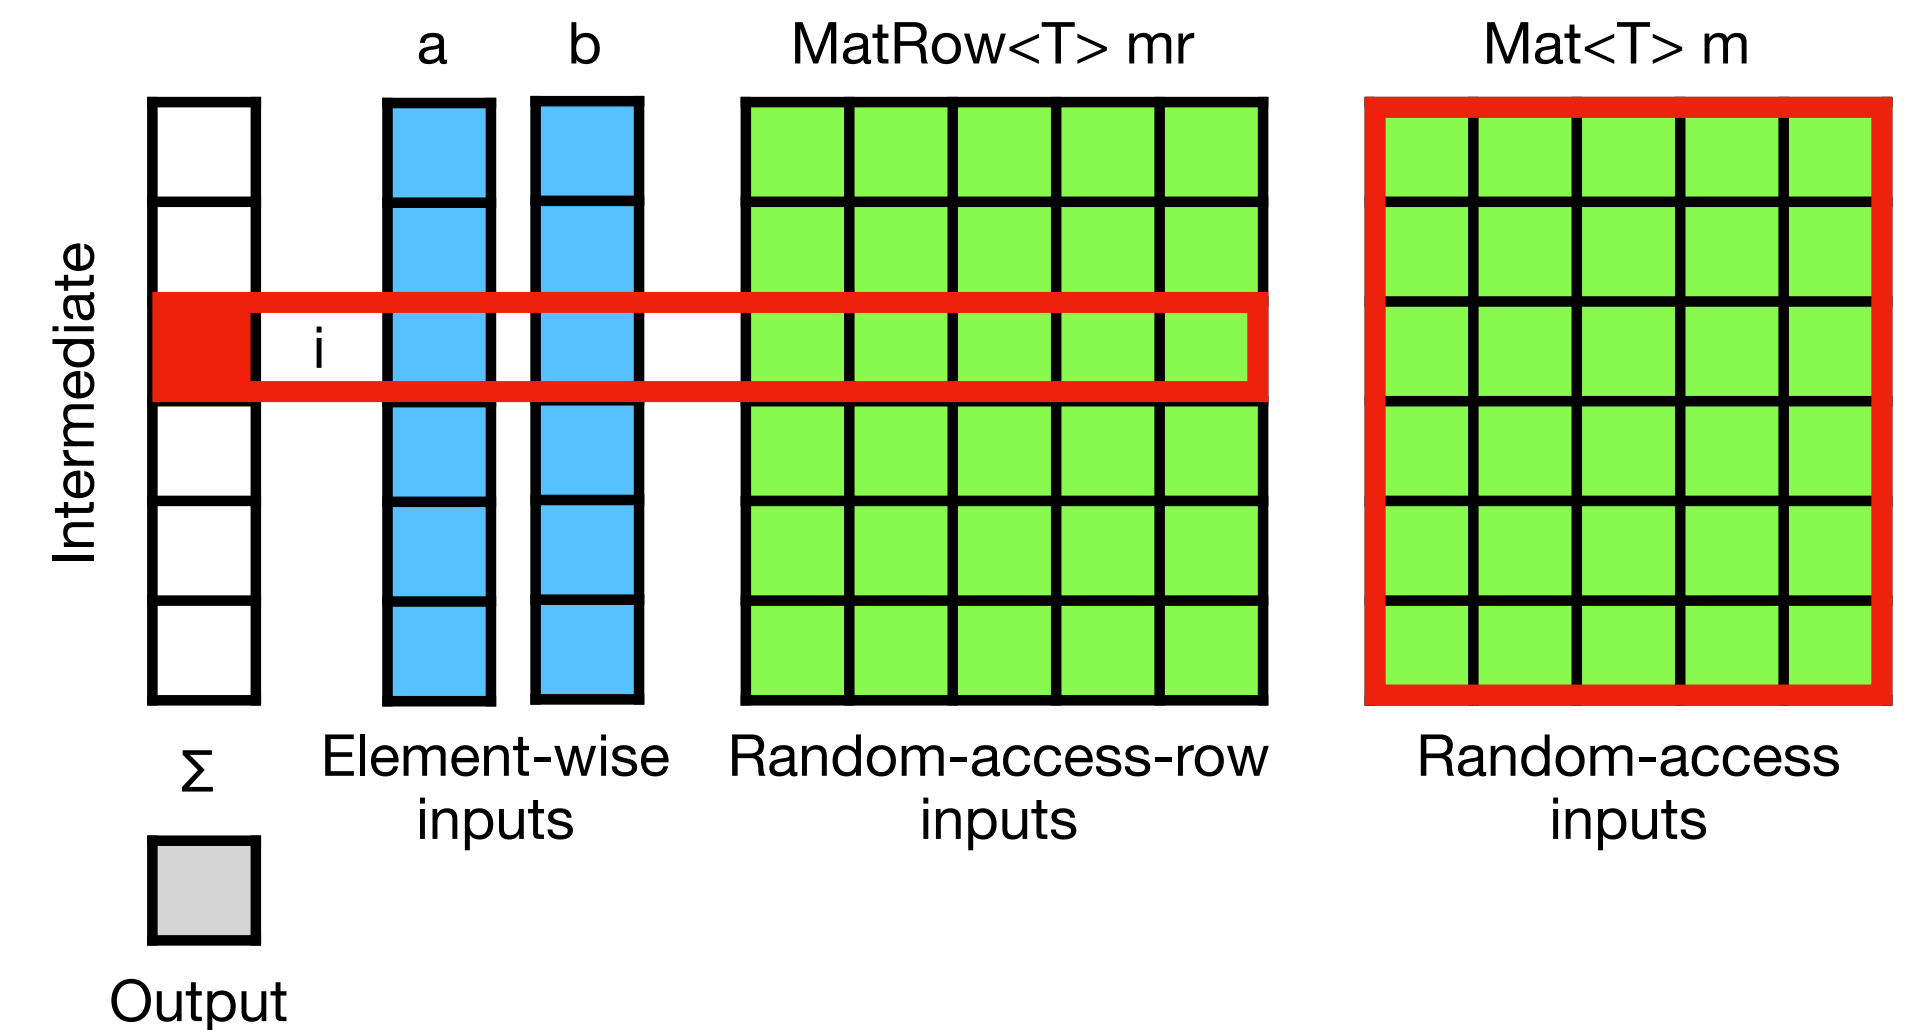
```
float max_f(float a, float b)
{
    return (a > b) ? a : b;
}
```

```
auto max_sum = skepu::Reduce(plus_f, max_f);
```

```
max_sum.setReduceMode(skepu::ReduceMode::RowWise);
r = max_sum(m);
```

MapReduce

- `instance.setSize(size_t)`
 - When the element-wise arity is 0, this controls the number of user function invocations (That is, the size of the “virtual” temporary container in between the Map and Reduce steps)
- `instance.setStartValue(value)`
 - Set Reduction start value. Defaults to 0-initialized.




```
float add(float a, float b)
{
    return a + b;
}
```

```
float mult(float a, float b)
{
    return a * b;
}
```

```
float dot_product(Vector<float> &v1, Vector<float> &v2)
{
    auto dotprod = MapReduce<2>(mult, add);
    return dotprod(v1, v2);
}
```

Scan

- `instance.setScanMode(mode)`
 - Set the scan mode:
`ScanMode::Inclusive` (default)
`ScanMode::Exclusive`
- `instance.setStartValue(value)`
 - Set start value of scans. Defaults to 0-initialized.

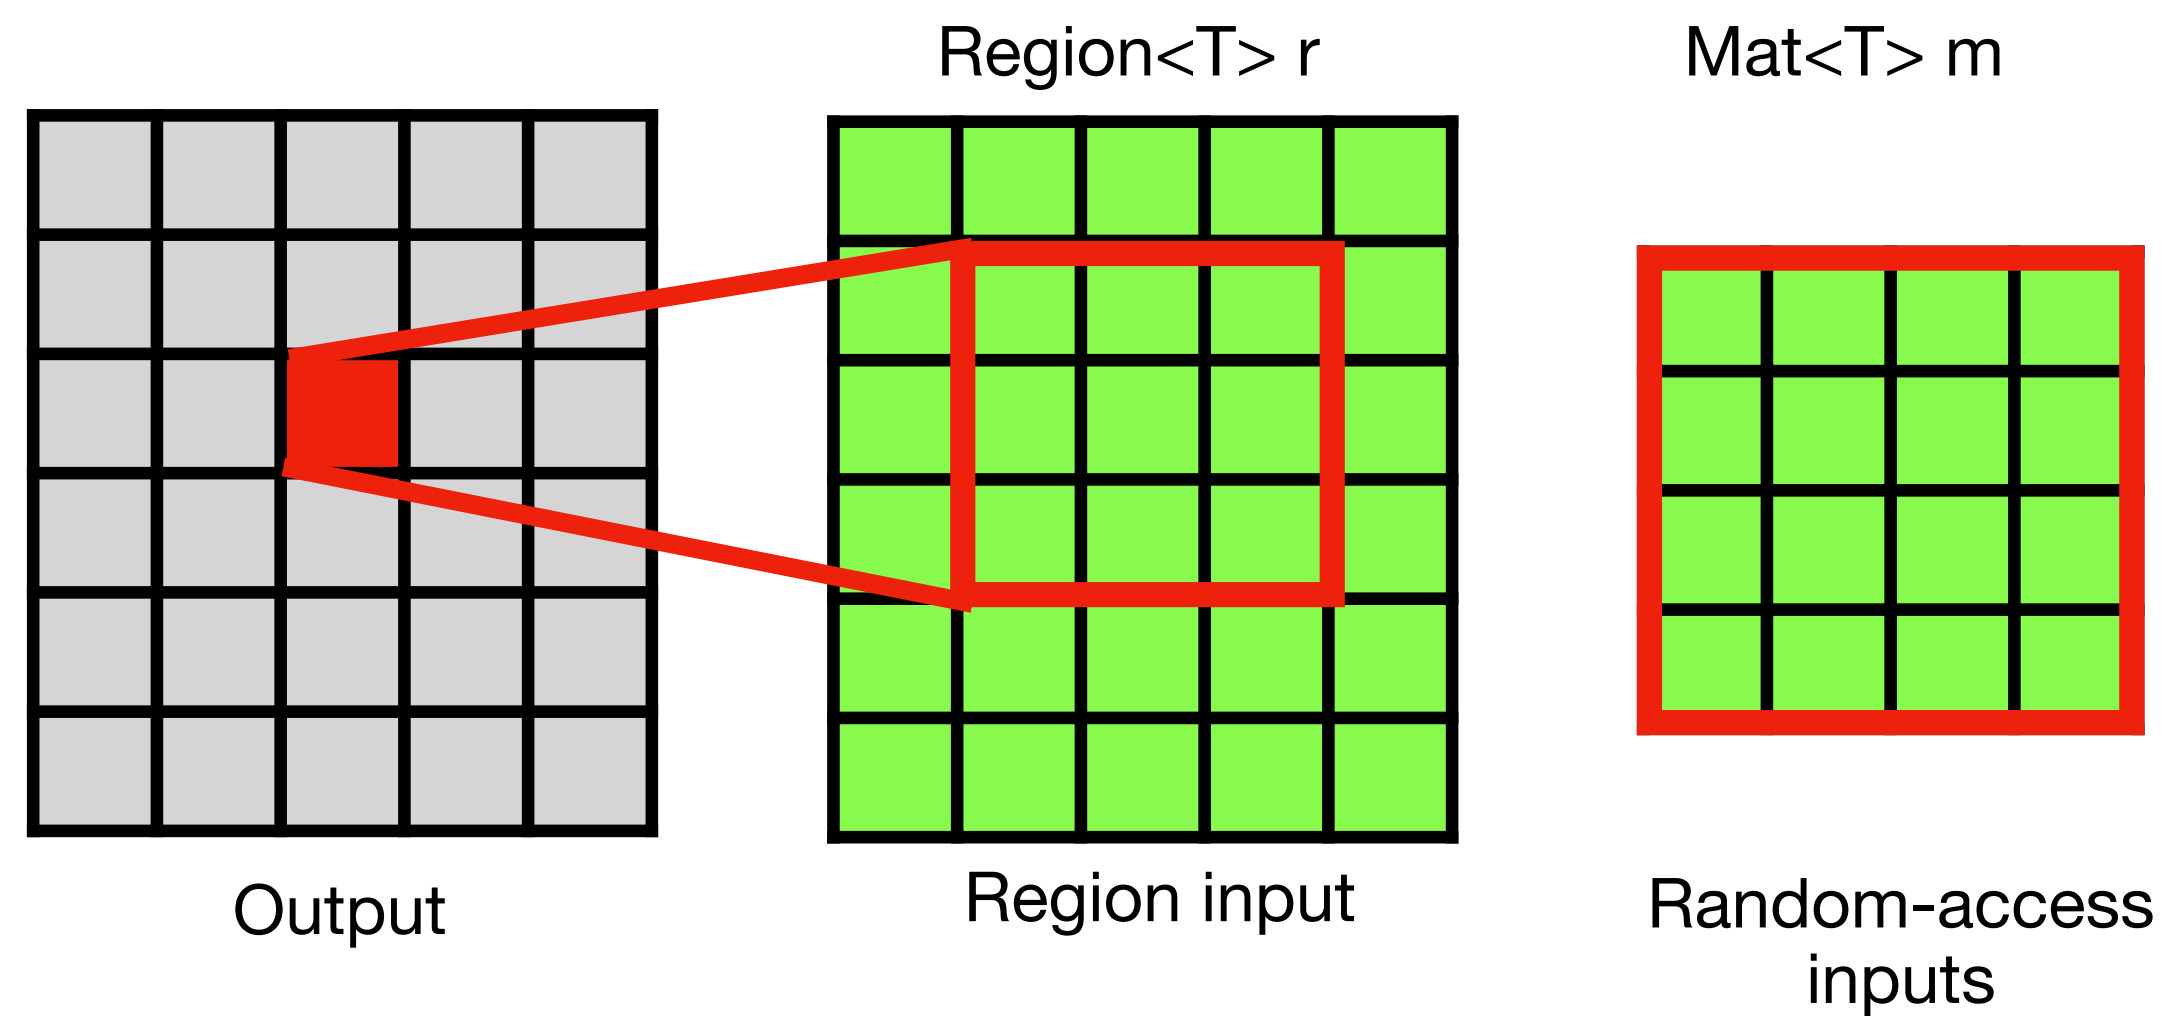


```
float max_f(float a, float b)
{
    return (a > b) ? a : b;
}
```


```
Vector<float> partial_max(Vector<float> &v)
{
    auto premax = Scan(max_f);
    Vector<float> result(v.size());
    return premax(result, v);
}
```

MapOverlap

- Region of the input container accessible in user function
- In addition to optional full random-access parameters



- Region indexing is 0-centered in each dimension





- **1D MapOverlap**

- Regular Vector



- Matrix RowWise `instance.setOverlapMode(Overlap::RowWise) // default`


- Matrix ColWise `instance.setOverlapMode(Overlap::ColWise)`



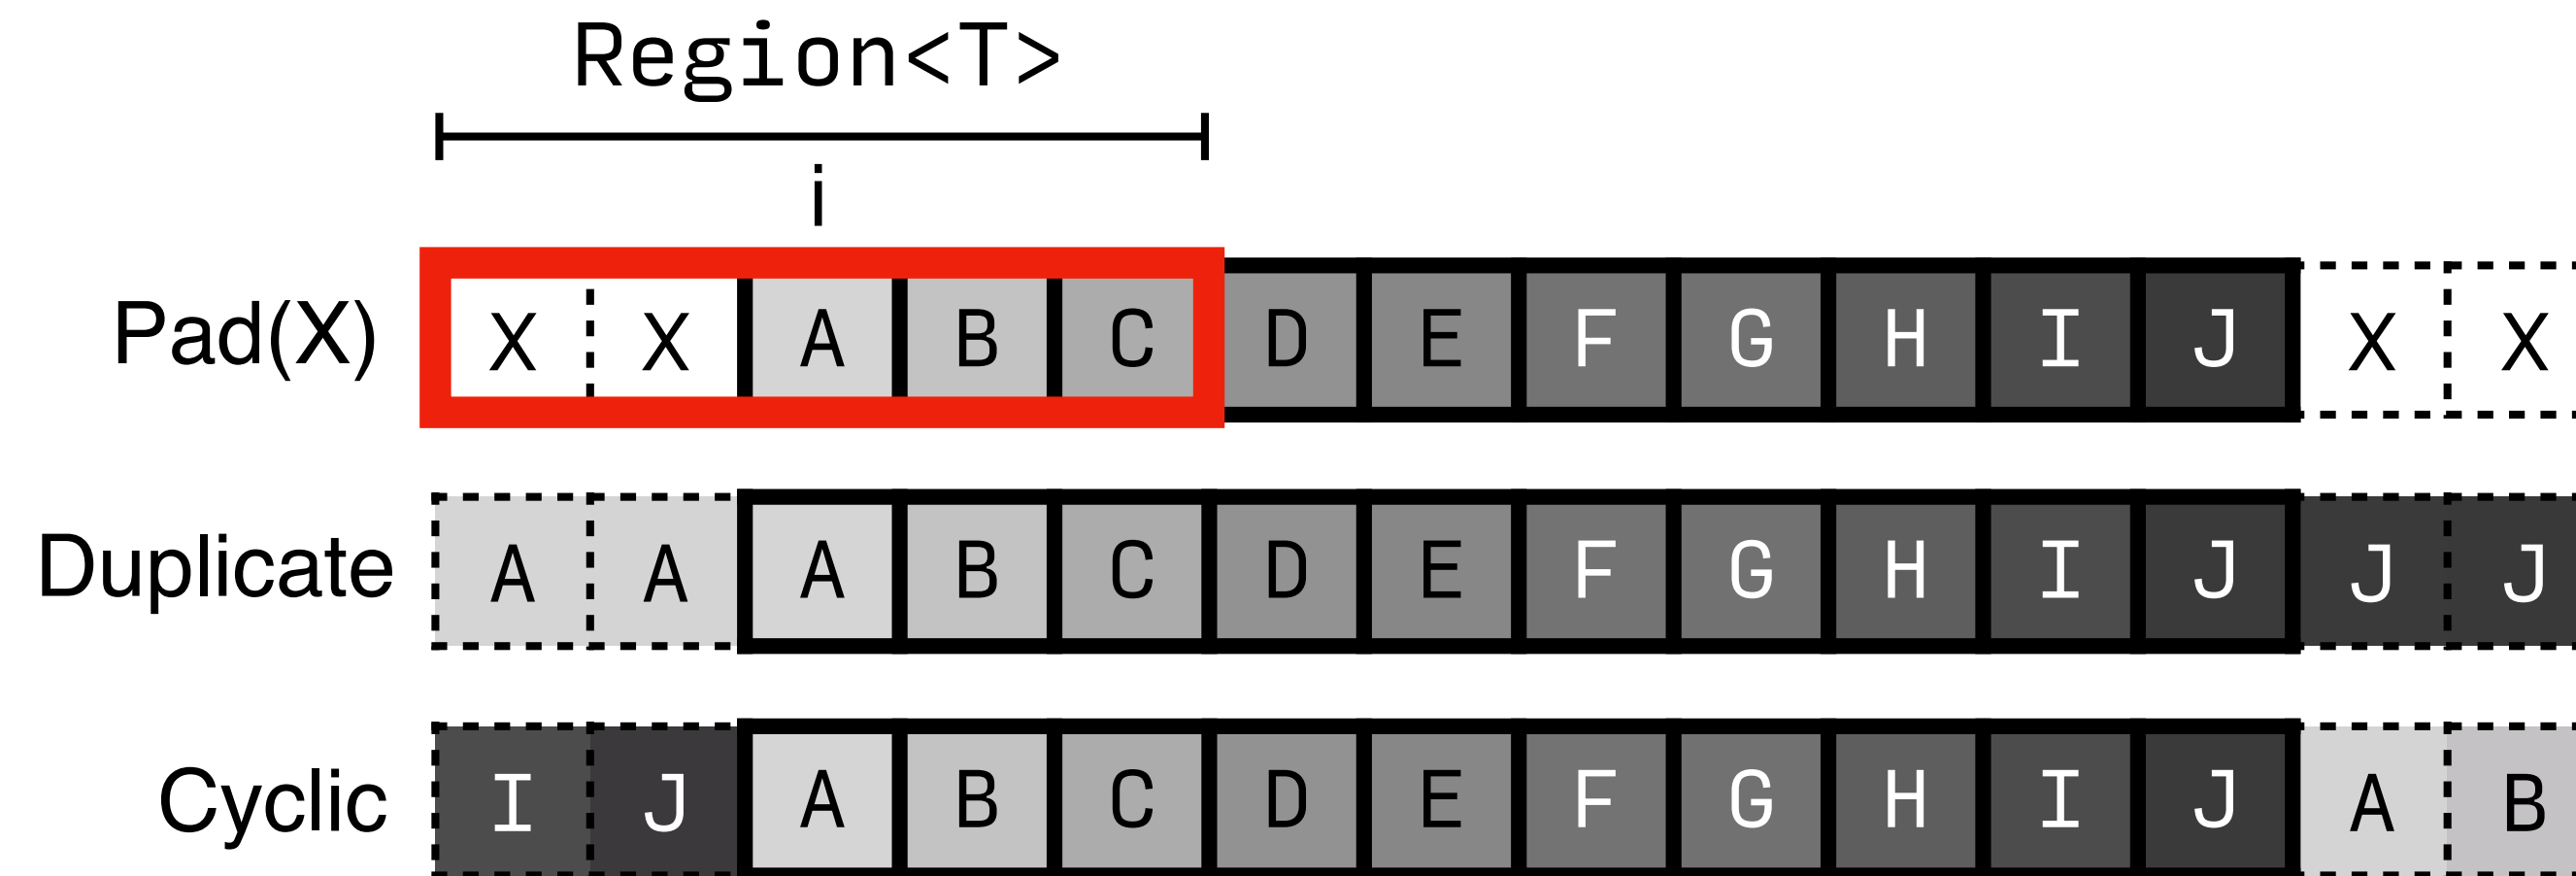
- **2D MapOverlap on Matrix**

- **3D MapOverlap on Tensor3**


- **4D MapOverlap on Tensor4**



- `instance.setOverlap(i [, j [, k [, 1]]])`
 - Set overlap radius in each dimension
- `instance.setEdgeMode(mode)`
 - `Edge::Pad`
 - `instance.setPad(pad)` – set padding value
 - `Edge::Duplicate` (default for 1D)
 - `Edge::Cyclic`
 - `Edge::None` (default for other dimensions)



None: Updates only "safe" elements




```
float conv(skepu::Region1D<float> r, int scale)
{
    return (r(-2)*4 + r(-1)*2 + r(0) + r(1)*2 + r(2)*4) / scale;
}

Vector<float> convolution(Vector<float> &v)
{
    auto convol = MapOverlap(conv);
    Vector<float> result(v.size());
    convol.setOverlap(2);
    return convol(result, v, 10);
}
```

```
float over_2d(skepu::Region2D<float> r, const skepu::Mat<float> stencil)
{
    float res = 0;
    for (int i = -r.oi; i <= r.oi; ++i)
        for (int j = -r.oj; j <= r.oj; ++j)
            res += r(i, j) * stencil(i + r.oi, j + r.oj);
    return res;
}
```

MapPairs

- Generalized cartesian product / outer product of vectors
- Always results in a Matrix
- Syntactically, Map extended with another variadic elwise parameter group




```


float ger_uf(float x, float y, float alpha)
{
    return alpha * x * y;
}

void outer_product(
    float alpha,
    Vector<float>& x,
    Vector<float>& y,
    Matrix<float>& A,
)
{
    auto skel = skepu::MapPairs<1, 1>(ger_uf);
    skel(A, x, y, alpha);
}

```

MapPairsReduce


- MapPairs chained with reduction along rows or columns of the matrix
- Efficient: matrix need not be allocated in full



Variadic Return

- Map*** skeletons optionally can return tuples of elements from the user function
- Compared to custom struct types, this stores results in several disjoint arrays

```
skel( resA, resB, inputs... );
```



```
skel( res, inputs... );
```



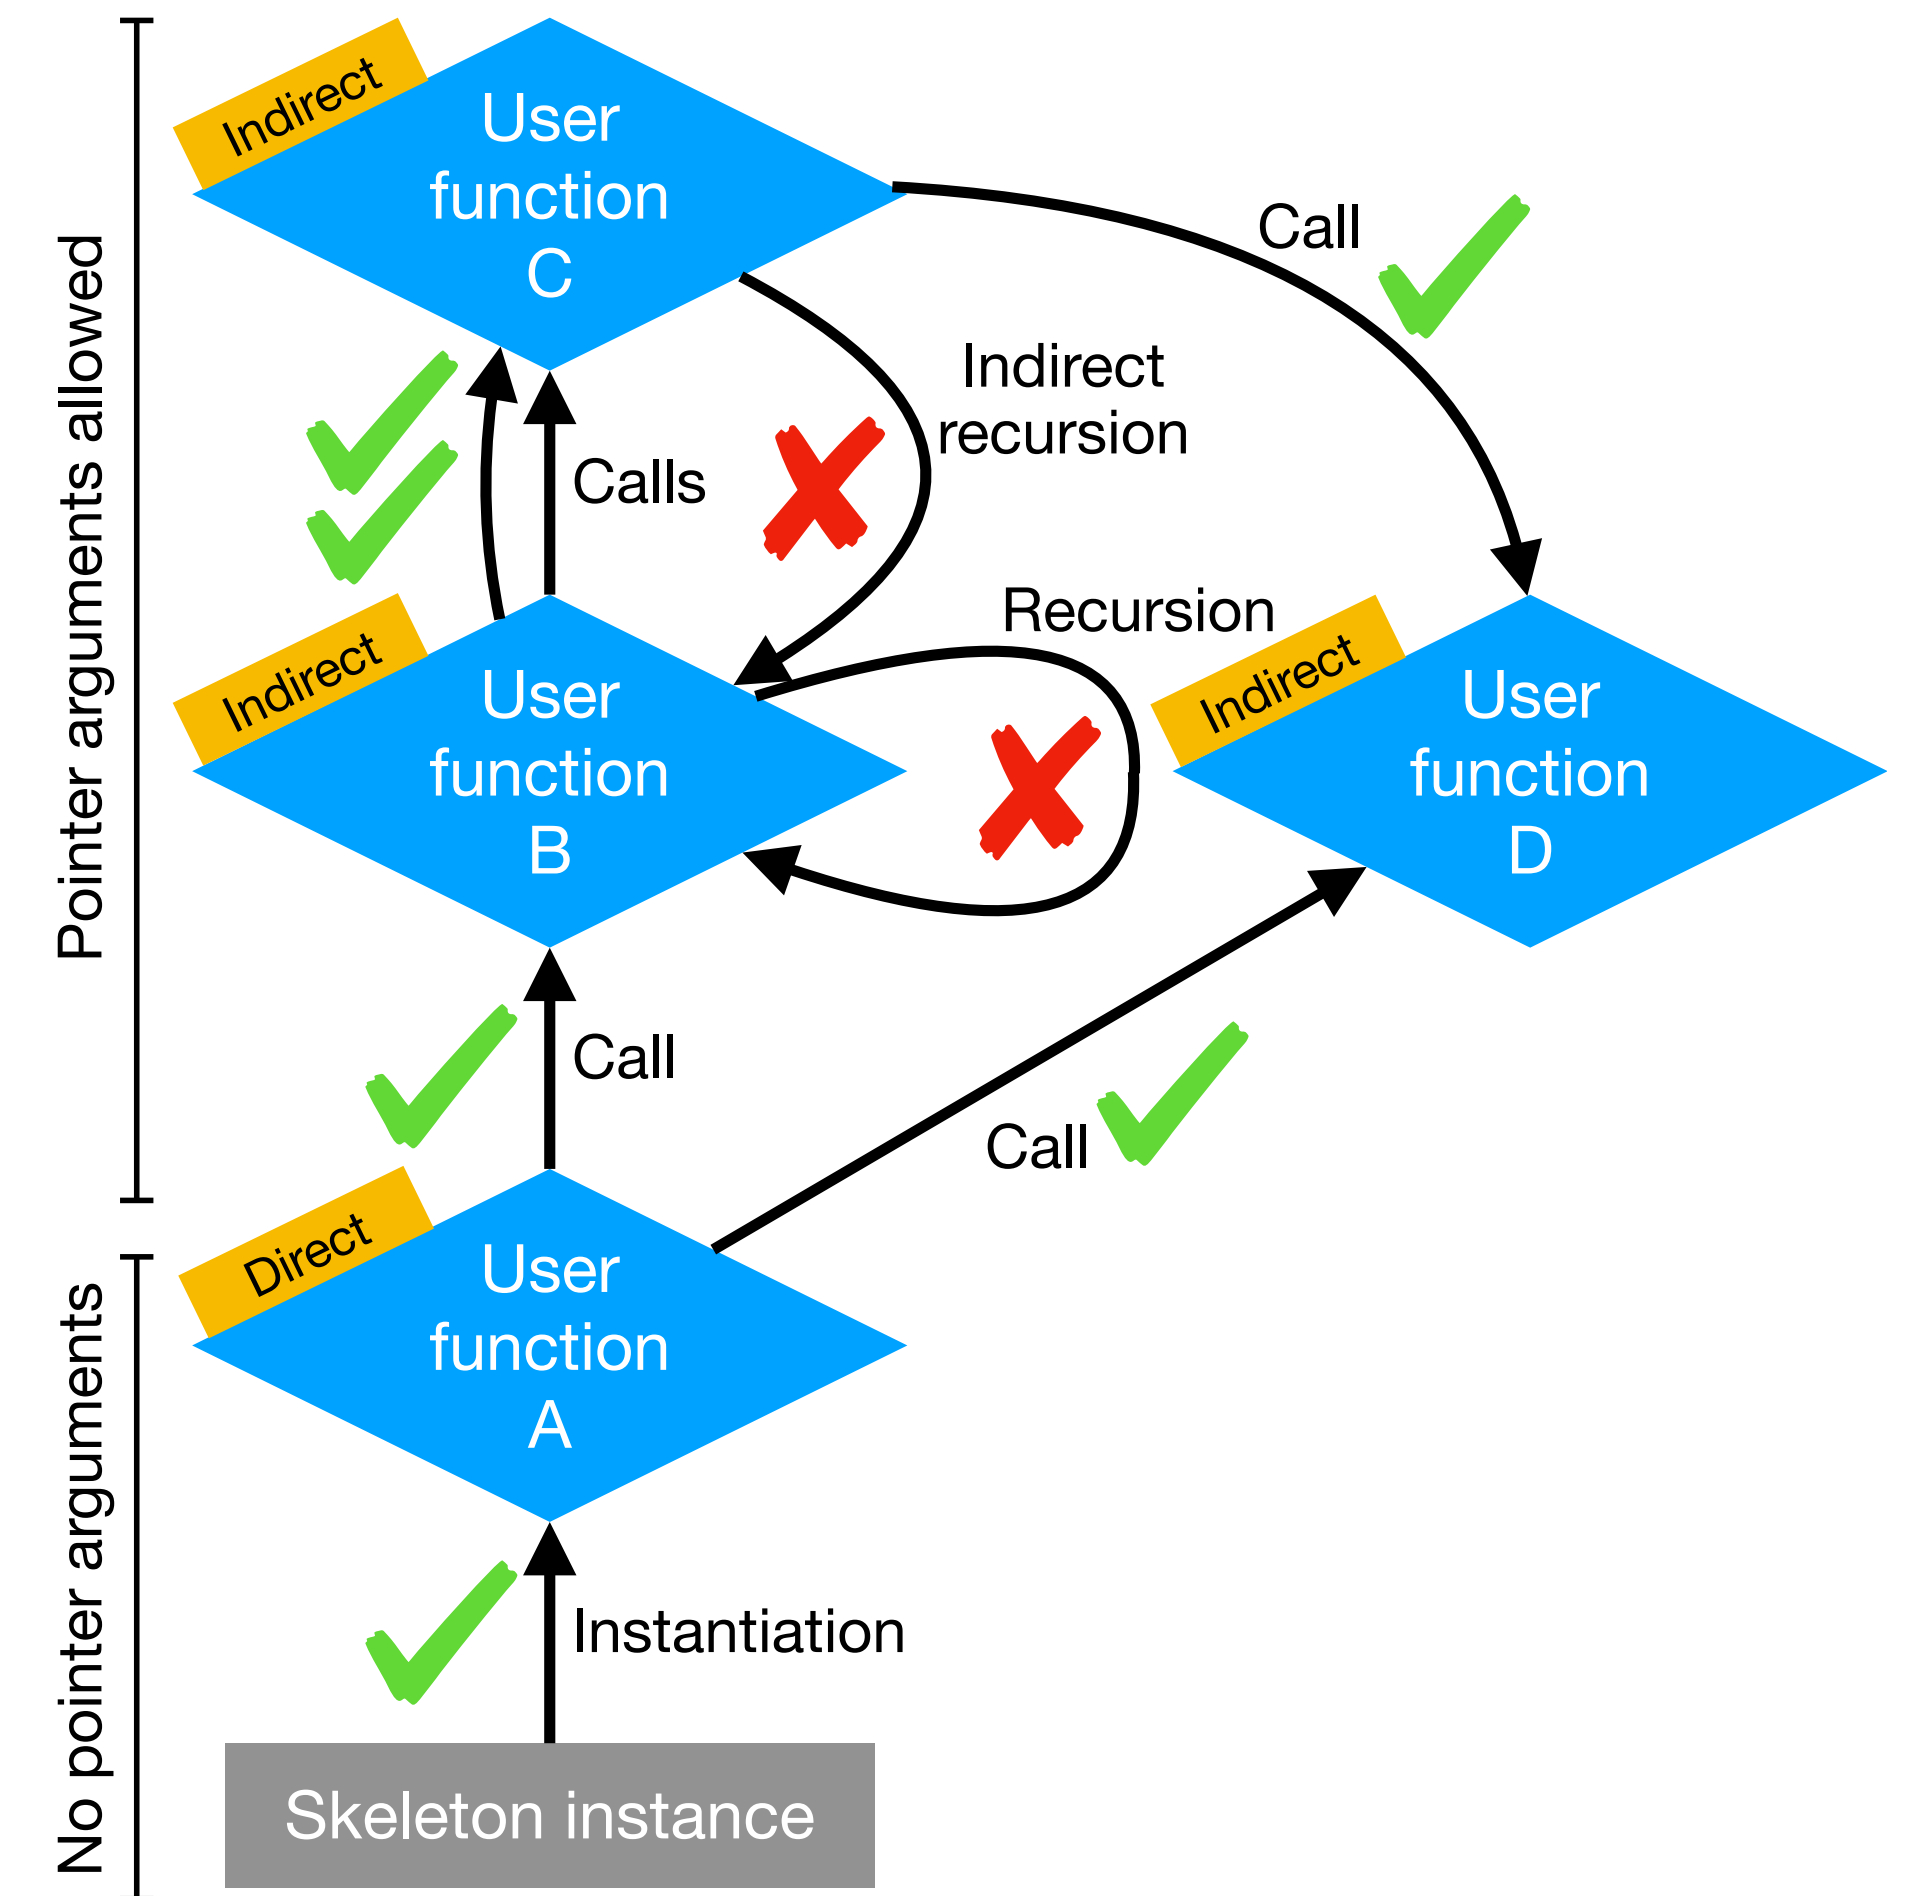
```
skepu::multiple<int, float>
test_f(skepu::Index1D index, int a, int b,
skepu::Vec<float> c, int d)
{
    return skepu::ret(a * b, (float)a / b);
}

// ...

auto test = skepu::Map<2>(test_f);
test(r1, r2, v1, v2, e, 10);
```

Nested User Functions

- User functions may call other functions
- The callee will be processed as its own user function by the precompiler
- Some restrictions are relaxed over the uf-uf barrier: pointers may be allowed



SkePU Standard Library


NB: Very new and subject to changes!

- SkePU "standard" library is a collection of **supporting APIs** to further simplify SkePU programming by eliminating the need for common boilerplate code.
- Some parts have framework integration, but most is implemented at **user-level** and can be used as reference for SkePU application programming.



skepu::Random<N>
skepu::PRNG

Deterministic PRNG




skepu::complex<float>
skepu::complex<double>

Complex number type,
functions, and operator
overloads




skepu::blas::axpy
skepu::blas::gemv
skepu::blas::gemm
...

Dense level 1-3 BLAS
implementation




skepu::filter::blur
skepu::filter::edge
skepu::filter::downsample
...

Image filtering routines
RGB + grayscale modes



skepu::io::cout
skepu::io::cin
skepu::io::cerr

I/O built on top of
skepu::external



Lambda expression-
based time
measurement

Statistical processing

Automated backend
alternation and
measurement



skepu::util::add
skepu::util::mul
skepu::util::identity
...

Common user functions

* Language/precompiler integration

- SkePU-accelerated BLAS implementation
 - User functions and multi-backend code generation “behind the scenes”
- Full coverage of level 1 BLAS
- Coverage of dense parts of level 2 + level 3
- Slightly modernized C++ function signatures
 - Otherwise argument-compatibility with CBLAS

● Conjugate gradient computation in SkePU-BLAS

```
#include <skepu>
#include <skepu-lib/blas.hpp>

template<typename T>
void conjugate_gradient(skepu::Matrix<T> BLAS_CONST& A, skepu::Vector<T> BLAS_CONST& b, skepu::Vector<T> &x)
{
    size_t N = b.size();
    assert(A.size_i() == N && A.size_j() == N && x.size() == N);
    skepu::Vector<T> p(N), r(N), Ap(N);

    // Set up initial r and p
    skepu::blas::copy(N, b, 1, r, 1);
    skepu::blas::gemv(skepu::blas::Op::NoTrans, N, N, -1.f, A, N, x, 1, 1.f, r, 1);
    skepu::blas::copy(N, r, 1, p, 1); // p := r

    float rTr = skepu::blas::dot(N, r, 1, r, 1); // rTr = r * r

    for (size_t k = 0; k < N; ++k)
    {
        // Compute alpha
        skepu::blas::gemv(skepu::blas::Op::NoTrans, N, N, 1.f, A, N, p, 1, 0.f, Ap, 1); // Ap := A * p
        float tmp = skepu::blas::dot(N, p, 1, Ap, 1); // tmp := p * Ap = p * A * p
        float alpha = rTr / tmp;

        // Update x
        skepu::blas::axpy(N, alpha, p, 1, x, 1); // x := x + alpha * p

        // Update r
        skepu::blas::axpy(N, -alpha, Ap, 1, r, 1); // r := r - alpha * Ap

        // Compute beta
        float rTr_new = skepu::blas::dot(N, r, 1, r, 1); // rTr_new := r * r
        float beta = rTr_new / rTr;

        // Early exit condition
        if (sqrt(rTr_new) < 1e-10f)
            return;

        // Update p
        skepu::blas::scal(N, beta, p, 1); // p := beta * p
        skepu::blas::axpy(N, 1.f, r, 1, p, 1); // p := r + p

        rTr = rTr_new;
    }
}
```

- `skepu::Random<N>` as user function parameter
 - Always first in parameter list
 - Maximum of one such parameter may be present
 - `.get()` extracts integers in $[0, \text{MAX})$
 - `.getNormalized()` extracts real numbers in $[0, 1)$
- `skepu::PRNG`
 - Outside user functions
 - Represents individual PRNG streams to bind to skeleton instances

● Monte-Carlo π calculation

```
#include <iostream>
#include <skepu>

int monte_carlo_sample(skepu::Random<2> &random)
{
    float x = random.getNormalized();
    float y = random.getNormalized();
    // check if (x,y) is inside region:
    return ((x*x + y*y) < 1) ? 1 : 0;
}

int add(int lhs, int rhs) { return lhs + rhs; }


int main(int argc, char *argv[])
{
    auto montecarlo = skepu::MapReduce<0>(monte_carlo_sample, add);

    skepu::PRNG prng(0); // seed
    montecarlo.setPRNG(prng);

    const size_t samples = atoi(argv[1]);
    montecarlo.setDefaultSize(samples);

    double pi = (double)montecarlo() / samples * 4;
    std::cout << pi << "\n";
}
```


- SkePU is well-suited for image processing with its **Map** and **MapOverlap** skeletons and **GPU** backends
- Image processing is also useful for **visualizing** unrelated HPC applications, e.g. heat diffusion
- Standard library includes **image filtering** primitives
 - Data types for **pixels** in various formats
 - Grayscale, RGB, RGBA, HSV, HSVA
 - **Filter kernel** user functions, e.g.
 - Pixel conversions
 - Channel adjustments (brightness, contrast, ...)
 - Convolutions (blurs, edge detection, ...)



```
#include <skepu>
#include <skepu-lib/filter.hpp>

skepu::Matrix<skepu::filter::RGBPixel> image;
// (Load image from file here)

auto lighten  = skepu::Map<1>(skepu::filter::lighten_rgb);
auto saturate = skepu::Map<1>(skepu::filter::saturate_rgb);

lighten(image, image, 0.5); // In-place update
saturate(image, image, 0.5); // In-place update
```

- SkePU standard library includes frequently seen user functions in the `skepu::util` namespace
- Arithmetic operations
- Identity mapping, type conversions
- Max, min, ...

```
#include <skepu>
#include <skepu-lib/util.hpp>

skepu::Vector<float> a(N), b(N);

auto dotproduct = skepu::MapReduce(skepu::util::mul<float>, skepu::util::add<float>);

float res = dotproduct(a, b);
```

SkePU in **Current Research**

Lazy Skeleton Evaluation

Ernstsson, Kessler, "Extending smart containers for data locality-aware skeleton programming",
Concurrency and Computation Practice and Experience, 2019; 31:e5003 <https://doi.org/10.1002/cpe.5003>

- SkePU extended with **lazy evaluation** of the Map skeleton
- No change in syntax, lineage graph built automatically at run-time
- Existing smart containers extended with some of the RDD semantics
- Non-Map operations cause lineage to be evaluated
 - Possibly out of order, while satisfying dependencies
- Run-time dataflow information can assist backend selection

```

add(v1, v3, v4);

copy(v9, v1);

mult(v2, v1, v3);

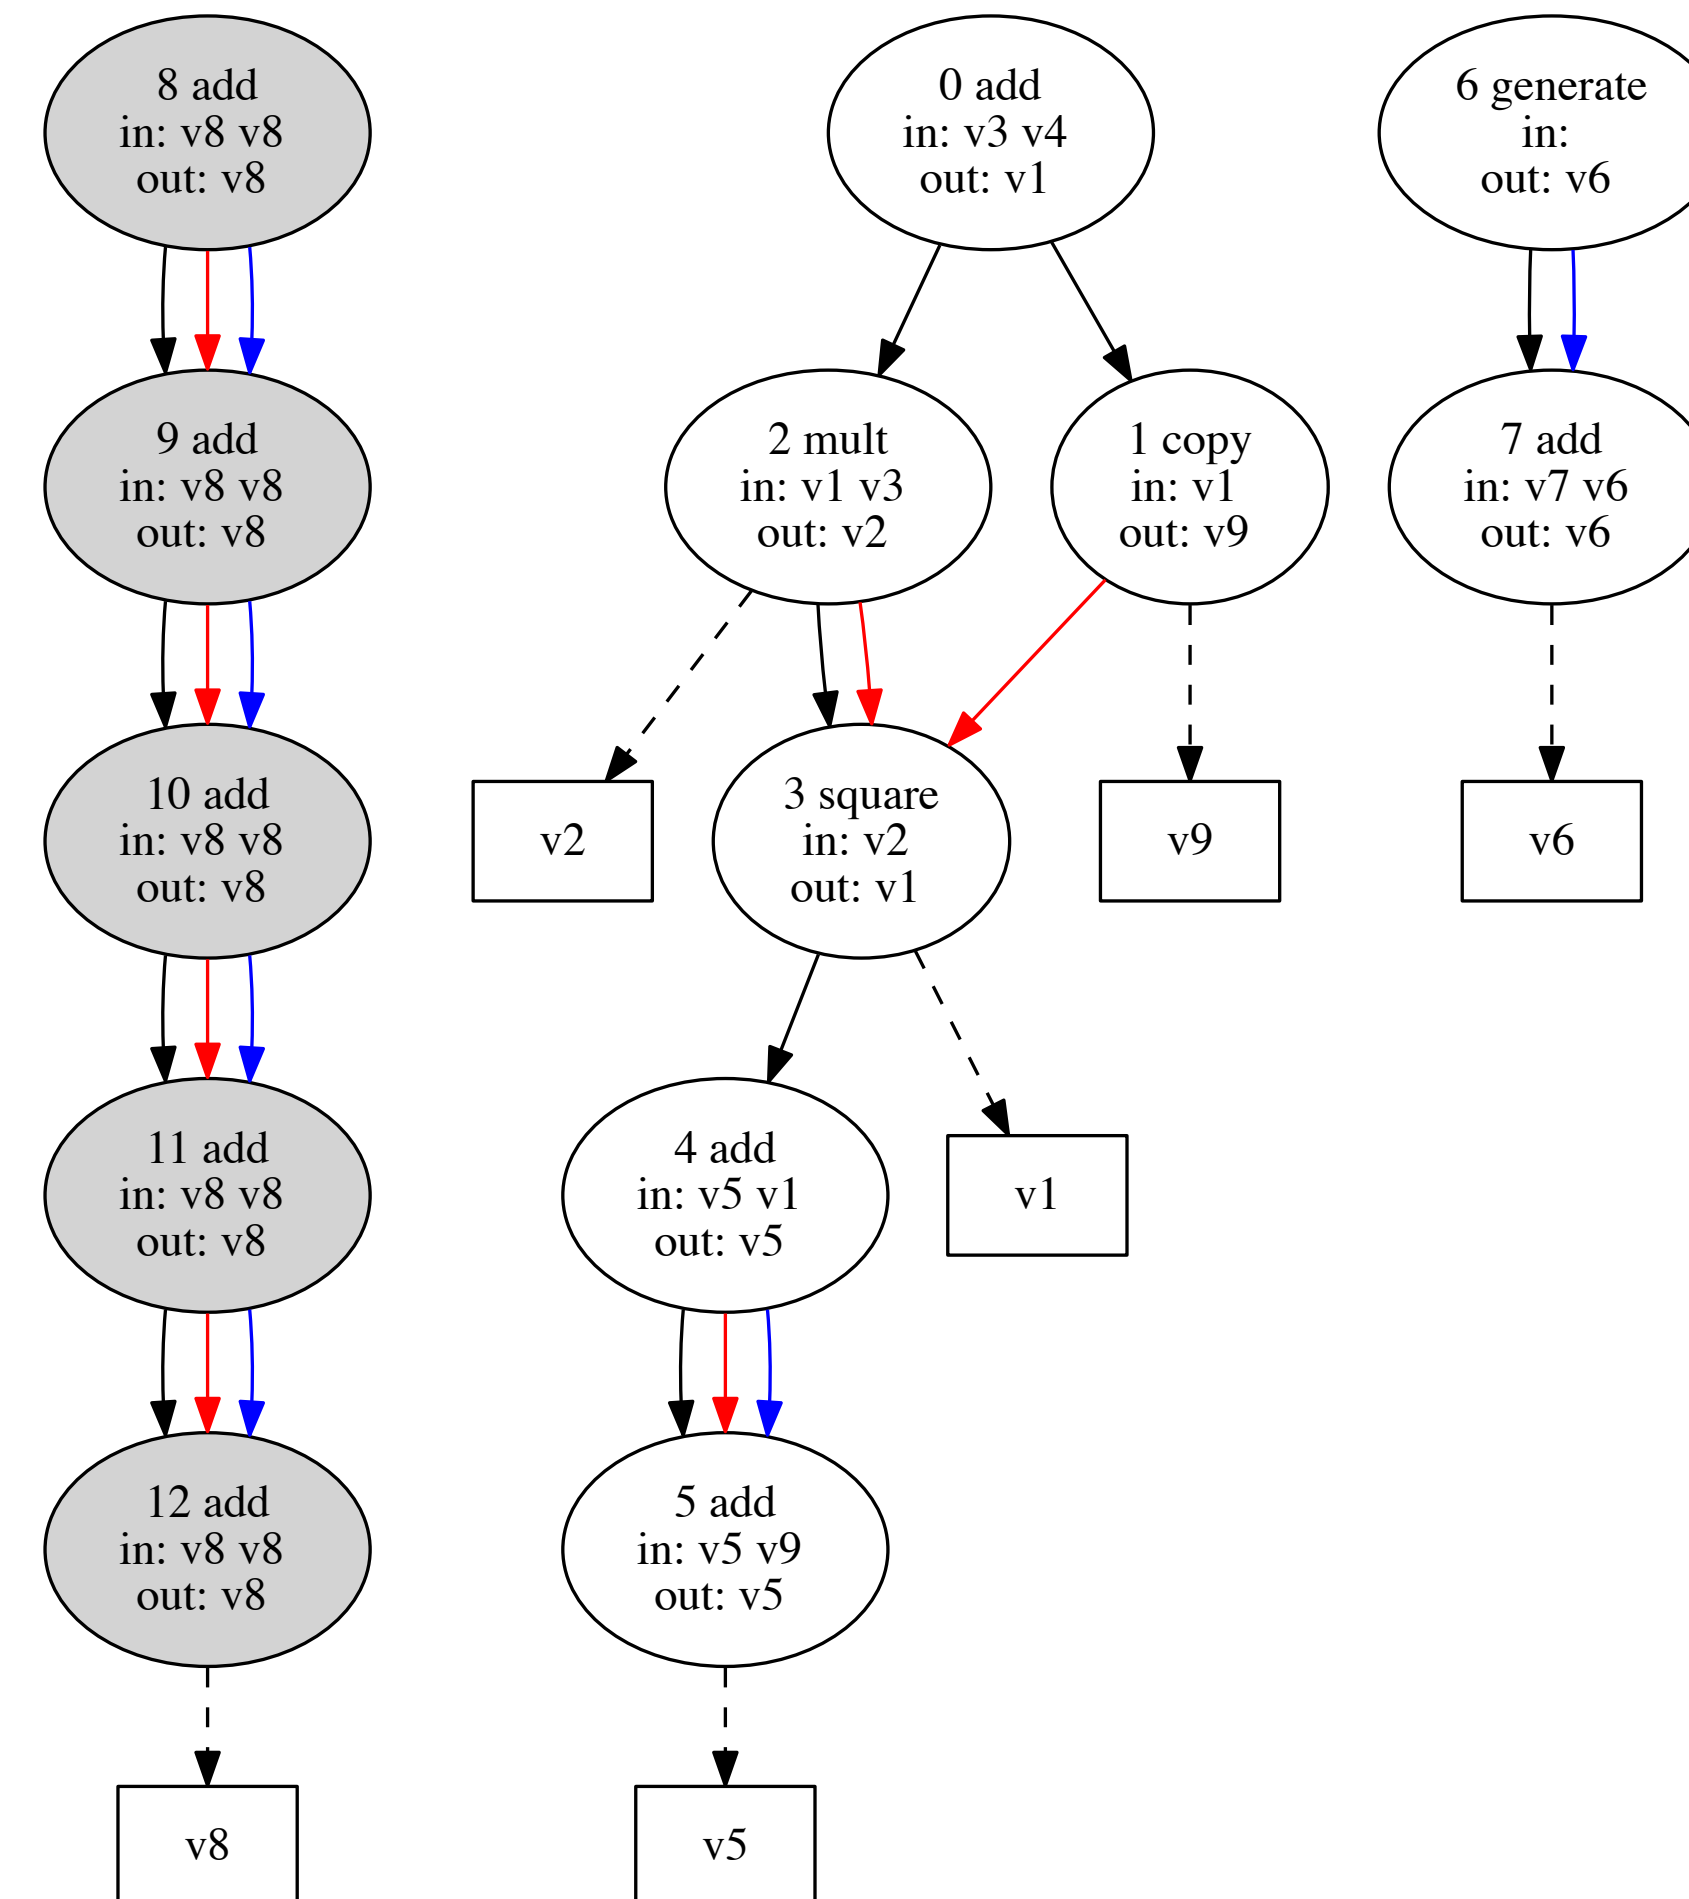
square(v1, v2);

add(v5, v5, v1);

add(v5, v5, v9);

add(v6, v7, generate(v6, 5.f));

for (int i = 0; i < 5; i++)
  add(v8, v8, v8);
  
```



SkePU lineage graph from program

Compile-time


SkePU precompiler
with transformations
Static dataflow information

Run-time


Auto-tuning
Smart containers
Dynamic dataflow information

- Lineages evaluated by **tiling** computations
 - Efficient cache utilization on CPU
 - Using accelerators with limited memory space
 - Any situation where cached data is important for performance

Standard SkePU



Lazy SkePU

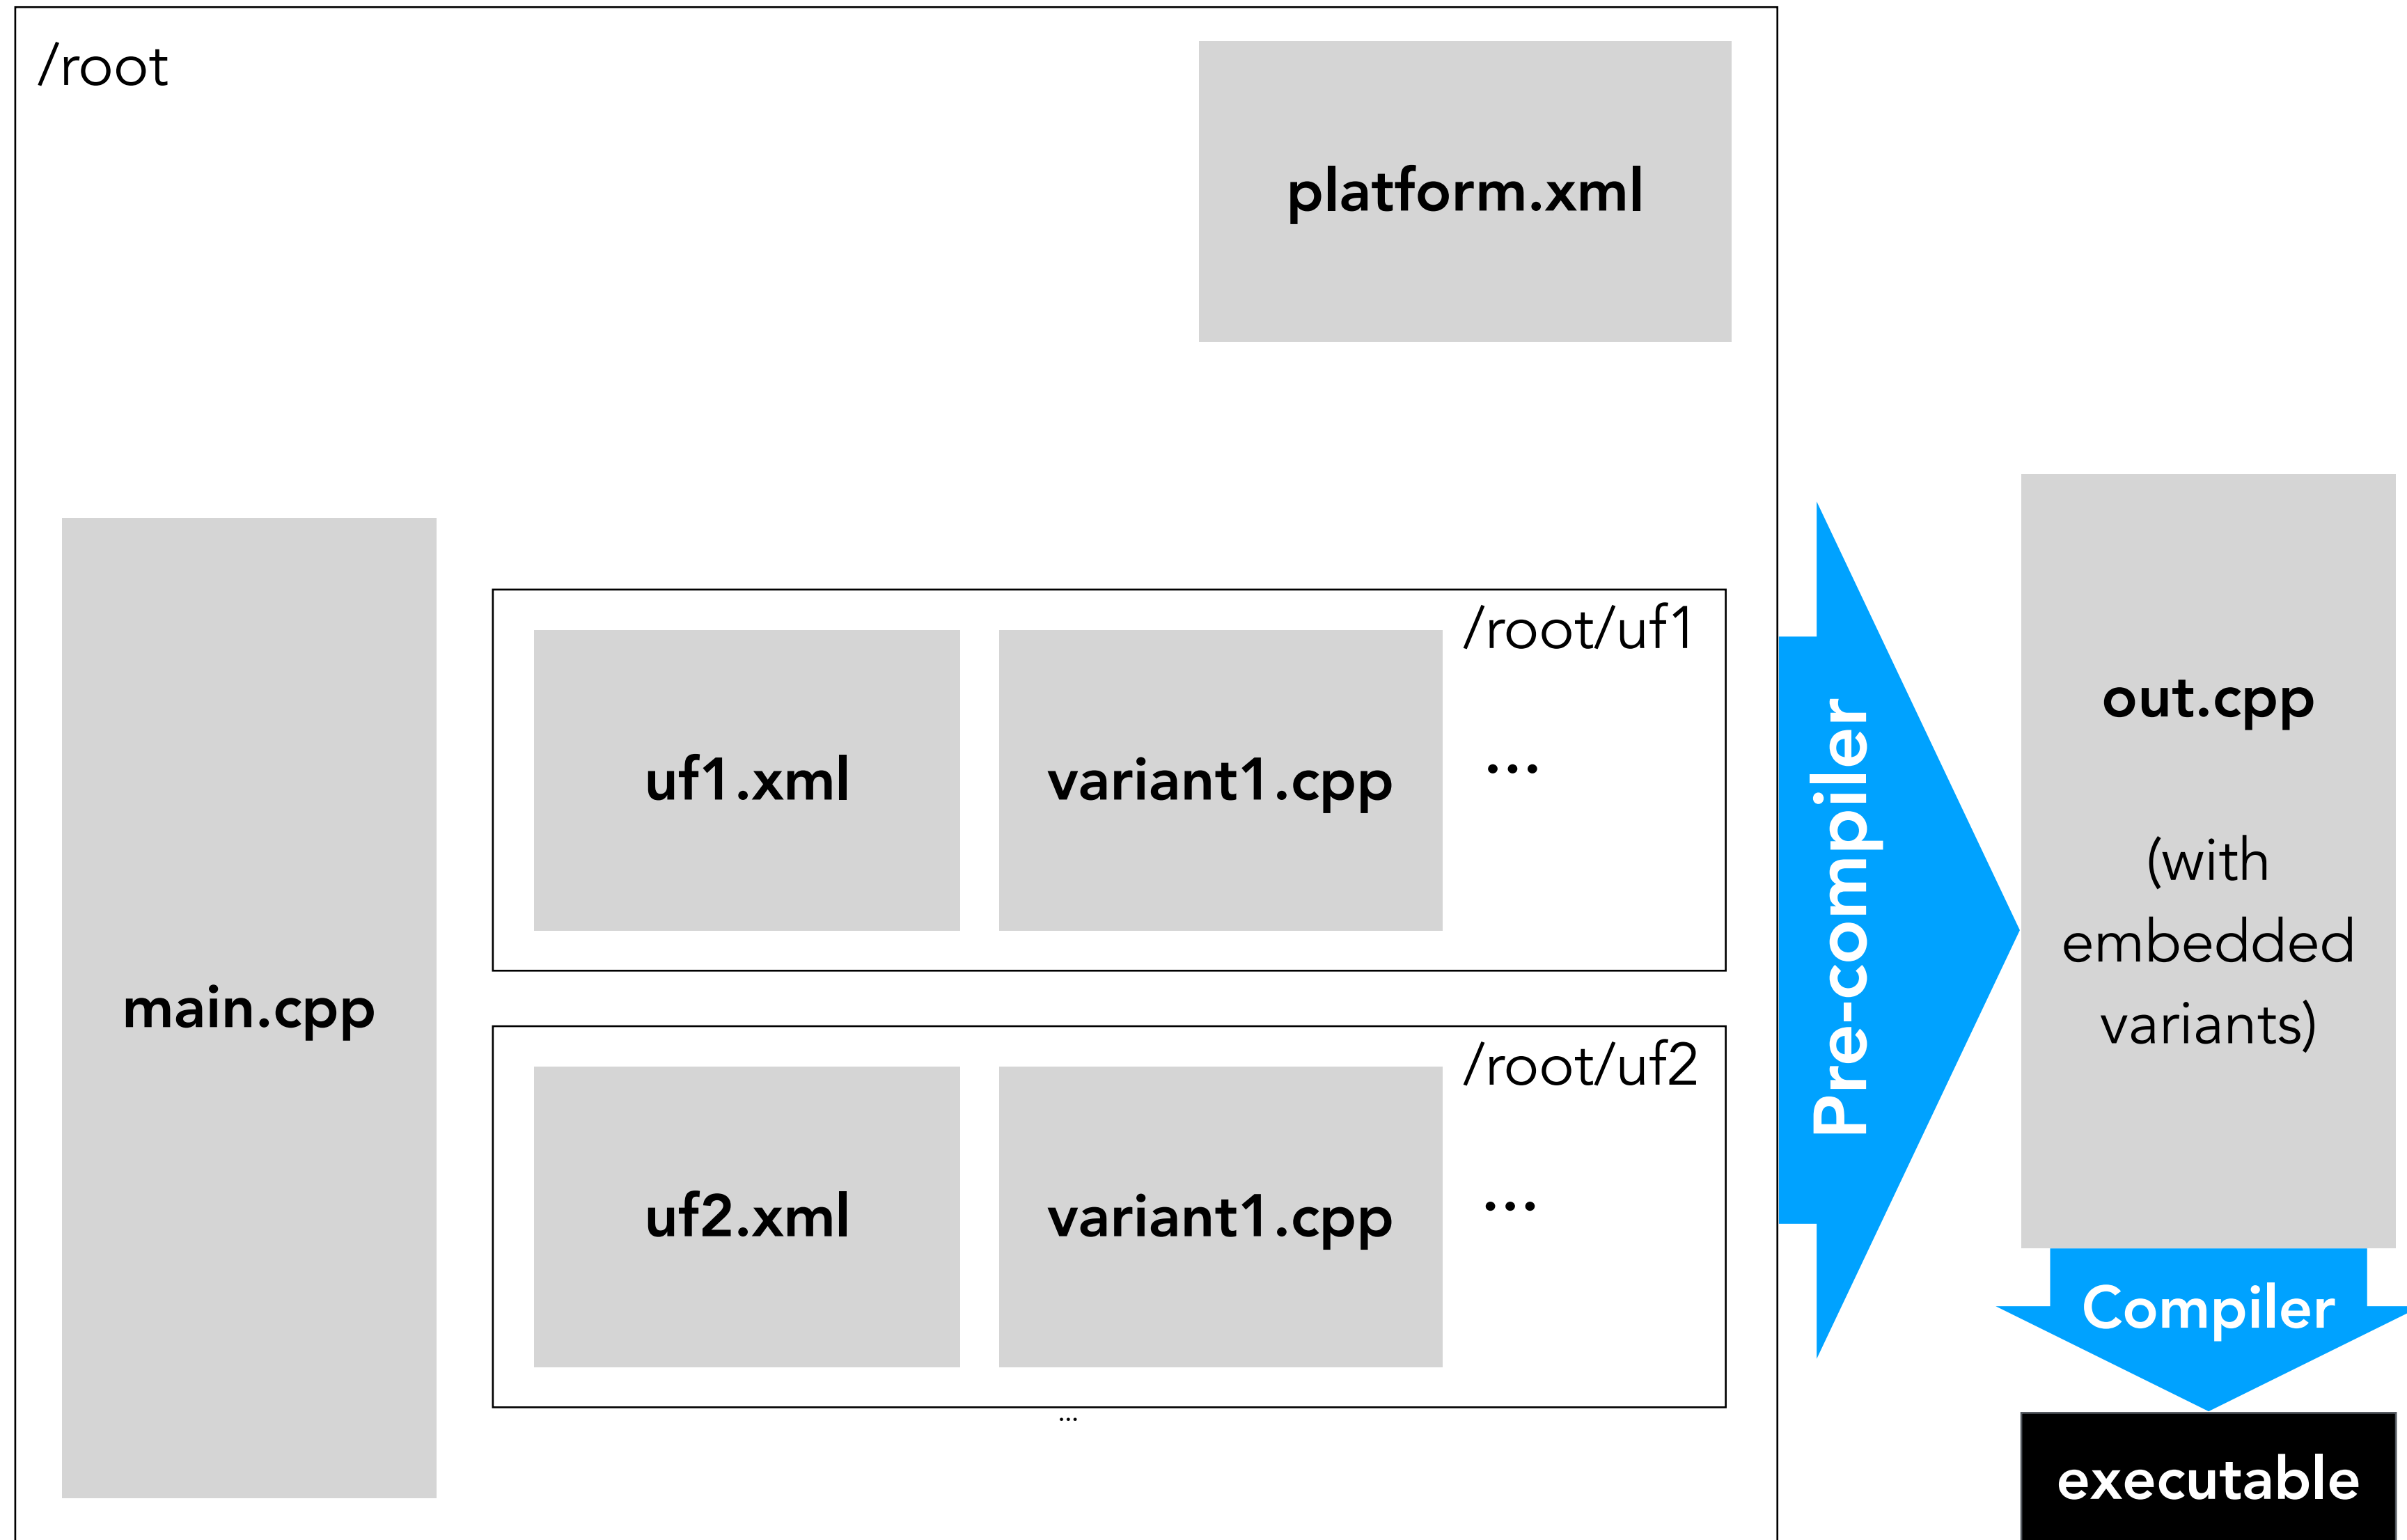


Multi-Variant User Functions




Ernstsson, Kessler, "Multi-Variant User Functions for Platform-Aware Skeleton Programming"
Parallel Computing: Technology Trends, IOS PRESS , 2020, Vol. 36, pp. 475-484
<http://doi.org/10.3233/APC200074>


- Single-source, in-line user function definitions are **easy to use**
- **But:** might not always be **performance optimal** approach across backends!
- **Idea:** Leverage multi-variant component approach in Exa2Pro also on skeleton/SkePU level
- Let each user function have multiple *implementation* variants conforming to the same *interface*
 - Algorithmic variants
 - Hardware/ISA optimizations (vectorize, crypto, ...)
 - External library acceleration, ...
- Each variant only selectable when the target system **supports** it
=> platform-aware (XPDL) source-to-source compilation



- Data-parallel algorithms using Map skeleton
- Archetype of skeleton programming: element-wise vector sum
- To increase the complexity, also do element-wise multiplication of complex numbers



- Median filtering: a blurring **image filter** with MapOverlap
- Data-parallel in the "outer" dimensions
- **Multiple algorithmic** approaches in the "inner" dimensions
 - Naive nested-loop comparison
 - Sorting
 - Histogram (note: data independent)




Variant	Time complexity	Memory complexity	Dependencies
Double loop	$\mathcal{O}(n^2)$	$\mathcal{O}(1)$	None
Histogram	$\mathcal{O}(n + D)$	$\mathcal{O}(D)$	None
qsort	$\mathcal{O}(n \log n)$	$\mathcal{O}(n)$	C standard library

Median Filtering Results

Variant	Time complexity	Memory complexity	Dependencies
Double loop	$\mathcal{O}(n^2)$	$\mathcal{O}(1)$	None
Histogram	$\mathcal{O}(n + D)$	$\mathcal{O}(D)$	None
qsort	$\mathcal{O}(n \log n)$	$\mathcal{O}(n)$	C standard library


■ Histogram@CPU ■ Double Loop@CPU ■ qsort@CPU
■ Histogram@OpenCL ■ Double Loop@OpenCL



Deterministic Pseudo- Random Number Generation



- Pseudo-random number generator (PRNG)
- **Finite** state automaton
 - `rand()` produces a pseudo-“random” number from its internal state $S(n)$ and transitions to new state $S(n+1)$
- Inherently **sequential**
- Naturally **deterministic**:
State $S(n)$ depends only on state $S(0)$ and n
- Input: **seed** (determines initial state)



```
void my_parallel_function(...)  
{  
    int rand_val = rand(); // library function  
}
```

Pseudocode

- **Race conditions**
 - Corrupted state
 - Repeated outputs
 - Loss of determinism



Bad statistical quality!

Global synchronization

```
void my_parallel_function(...)
{
    mutex_lock(...);
    int rand_val = rand(); // library function
    mutex_unlock(...);
}
```

Pseudocode

- Flexible

- Sequentialization
- Non-determinism
- Might not be supported (e.g. SkePU)
- Synchronization overhead

Pre-generated buffers

```
void my_parallel_function(int randvals[], ...)
{
    int i = my_thread_id();
    int rand_val = randvals[i++];
}
```


Pseudocode

- Simple: works like any data buffer
- Widely supported

- Not really parallel
- Manual pre-allocation of numbers
- Data movement costs



... we can do better!

- Basic idea: Give each worker thread a **separate** independent PRNG state
- No synchronization needed
- PRNG generation in parallel – no bulk data movement
- Individual PRNGs are **independent**
 - Loss of original single-PRNG structure
 - Need p seeds
 - Hierarchical seeding: **master PRNG** producing seeds
- This was roughly the approach for the first SkePUized LQCD prototype



... we can go further!

- **Goals:**
 - Retain the advantages of "stream splitting" approach
 - Reclaim structure and determinism from sequential PRNG
- **Recall:** PRNG state automaton
 - We can **advance** the state N steps at any point!
 - As long as we know #calls to PRNG in each thread



- Forwarding step incurs some overhead
- For **sequences** of skeleton calls, including **iterative** computations, we want to forward as few times as possible

```
for (int i = 0; i < N; ++i)
{
    my_skeleton_instance(...);
}
```

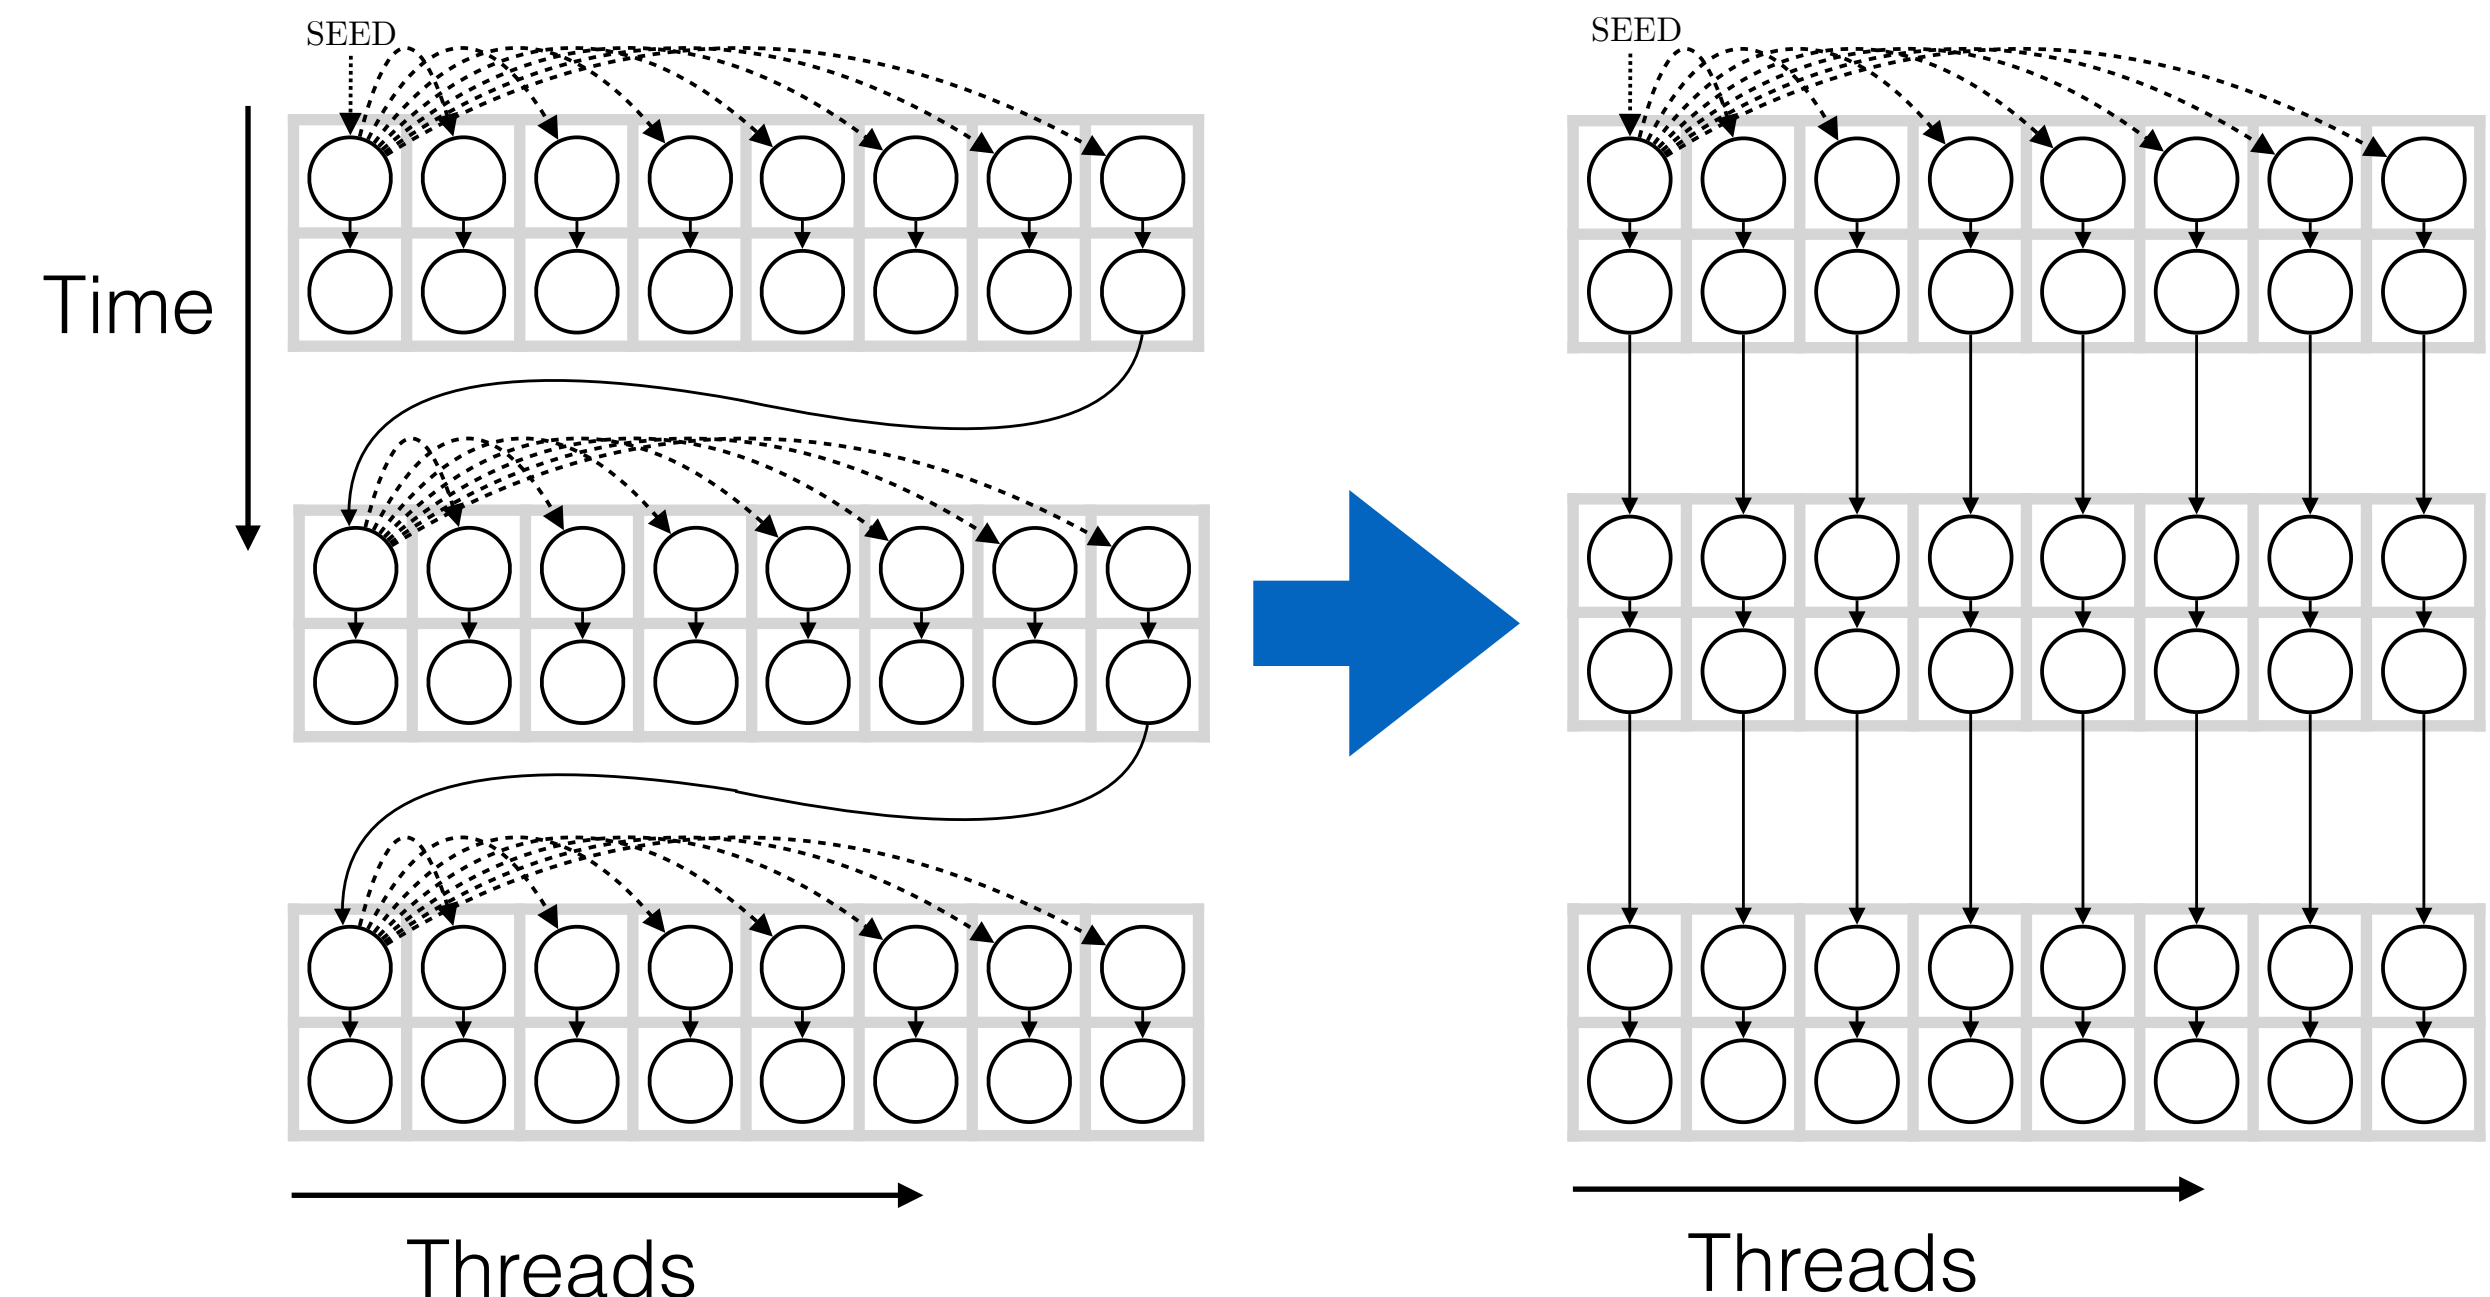
- **Pre-forward approach** for sequences of skeletons when program flow is known ahead of execution-time

- Static analysis

- Lineage building

Ernstsson, Kessler, HLPP 2017
<https://doi.org/10.1002/cpe.5003>

- Programmer annotation



Behind the Scenes

- SkePU is now **open-source on GitHub**: skepu.github.io
- Permissive modified BSD-license
- Systematic automated testing
- Continuously improving user guide, tutorials, etc.
- SkePU is a key framework in the H2020 EXA2PRO research project
- SkePU is actively used in **teaching** (parallel programming courses at LiU)

